

Importing Necessary Packages

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df=pd.read_csv(r"C:\Users\Arigala.Adarsh\Downloads\Data-sets\census-
income (7).csv")
```

```
df.head()
```

	age	workclass	fnlwgt	education	education-num \
0	39	State-gov	77516	Bachelors	13
1	50	Self-emp-not-inc	83311	Bachelors	13
2	38	Private	215646	HS-grad	9
3	53	Private	234721	11th	7
4	28	Private	338409	Bachelors	13

	marital-status	occupation	relationship	race
sex \				
0	Never-married	Adm-clerical	Not-in-family	White
Male				
1	Married-civ-spouse	Exec-managerial	Husband	White
Male				
2	Divorced	Handlers-cleaners	Not-in-family	White
Male				
3	Married-civ-spouse	Handlers-cleaners	Husband	Black
Male				
4	Married-civ-spouse	Prof-specialty	Wife	Black
Female				

	capital-gain	capital-loss	hours-per-week	native-country
0	2174	0	40	United-States
<=50K				
1	0	0	13	United-States
<=50K				
2	0	0	40	United-States
<=50K				
3	0	0	40	United-States
<=50K				
4	0	0	40	Cuba
<=50K				

Data Preprocessing

```
df.rename(columns={' ': 'Annual-Income'}, inplace=True)
```

```
df.head()
```

	age	workclass	fnlwgt	education	education-num \
0	39	State-gov	77516	Bachelors	13
1	50	Self-emp-not-inc	83311	Bachelors	13
2	38	Private	215646	HS-grad	9
3	53	Private	234721	11th	7
4	28	Private	338409	Bachelors	13

	marital-status	occupation	relationship	race
sex \				
0	Never-married	Adm-clerical	Not-in-family	White
Male				
1	Married-civ-spouse	Exec-managerial	Husband	White
Male				
2	Divorced	Handlers-cleaners	Not-in-family	White
Male				
3	Married-civ-spouse	Handlers-cleaners	Husband	Black
Male				
4	Married-civ-spouse	Prof-specialty	Wife	Black
Female				

	capital-gain	capital-loss	hours-per-week	native-country
Annuual-Income				
0	2174	0	40	United-States
<=50K				
1	0	0	13	United-States
<=50K				
2	0	0	40	United-States
<=50K				
3	0	0	40	United-States
<=50K				
4	0	0	40	Cuba
<=50K				

```
df.shape
```

```
(32561, 15)
```

```
df.columns
```

```
Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',  
      'marital-status', 'occupation', 'relationship', 'race', 'sex',  
      'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',  
      'Annuual-Income'],  
      dtype='object')
```

```
df.dtypes
```

```
age                int64
workclass          object
fnlwgt            int64
education          object
education-num      int64
marital-status     object
occupation         object
relationship       object
race              object
sex               object
capital-gain       int64
capital-loss       int64
hours-per-week     int64
native-country     object
Annuual-Income     object
dtype: object
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
```

#	Column	Non-Null	Count	Dtype
0	age	32561	non-null	int64
1	workclass	32561	non-null	object
2	fnlwgt	32561	non-null	int64
3	education	32561	non-null	object
4	education-num	32561	non-null	int64
5	marital-status	32561	non-null	object
6	occupation	32561	non-null	object
7	relationship	32561	non-null	object
8	race	32561	non-null	object
9	sex	32561	non-null	object
10	capital-gain	32561	non-null	int64
11	capital-loss	32561	non-null	int64
12	hours-per-week	32561	non-null	int64
13	native-country	32561	non-null	object
14	Annuual-Income	32561	non-null	object

```
dtypes: int64(6), object(9)
```

```
memory usage: 3.7+ MB
```

```
df.isnull().sum()
```

```
age                0
workclass          0
fnlwgt            0
education          0
education-num      0
marital-status     0
```

```

occupation      0
relationship    0
race            0
sex            0
capital-gain    0
capital-loss    0
hours-per-week  0
native-country  0
Annuual-Income  0
dtype: int64

df.duplicated()

0      False
1      False
2      False
3      False
4      False
...
32556   False
32557   False
32558   False
32559   False
32560   False
Length: 32561, dtype: bool

df.duplicated().sum()

24

df.drop_duplicates(inplace=True)

df.shape

(32537, 15)

```

Exploratory Data analysis-EDA

```
df.describe()
```

	age	fnlwgt	education-num	capital-gain	\
count	32537.000000	3.253700e+04	32537.000000	32537.000000	
mean	38.585549	1.897808e+05	10.081815	1078.443741	
std	13.637984	1.055565e+05	2.571633	7387.957424	
min	17.000000	1.228500e+04	1.000000	0.000000	
25%	28.000000	1.178270e+05	9.000000	0.000000	
50%	37.000000	1.783560e+05	10.000000	0.000000	
75%	48.000000	2.369930e+05	12.000000	0.000000	
max	90.000000	1.484705e+06	16.000000	99999.000000	

	capital-loss	hours-per-week
count	32537.000000	32537.000000
mean	87.368227	40.440329
std	403.101833	12.346889
min	0.000000	1.000000
25%	0.000000	40.000000
50%	0.000000	40.000000
75%	0.000000	45.000000
max	4356.000000	99.000000

```
df.columns
```

```
Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',
       'marital-status', 'occupation', 'relationship', 'race', 'sex',
       'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',
       'Annuual-Income'],
      dtype='object')
```

```
df['occupation'].unique()
```

```
array(['Adm-clerical', 'Exec-managerial', 'Handlers-cleaners',
       'Prof-specialty', 'Other-service', 'Sales', 'Craft-repair',
       'Transport-moving', 'Farming-fishing', 'Machine-op-inspct',
       'Tech-support', '?', 'Protective-serv', 'Armed-Forces',
       'Priv-house-serv'], dtype=object)
```

```
df['occupation'].nunique()
```

```
15
```

```
df['native-country'].unique()
```

```
array(['United-States', 'Cuba', 'Jamaica', 'India', '?', 'Mexico',
       'South', 'Puerto-Rico', 'Honduras', 'England', 'Canada',
       'Germany', 'Iran', 'Philippines', 'Italy', 'Poland',
       'Columbia', 'Cambodia', 'Thailand', 'Ecuador', 'Laos',
       'Taiwan', 'Haiti', 'Portugal', 'Dominican-Republic',
       'El-Salvador', 'France', 'Guatemala', 'China', 'Japan',
       'Yugoslavia', 'Peru', 'Outlying-US(Guam-USVI-etc)', 'Scotland',
       'Trinidad&Tobago', 'Greece', 'Nicaragua', 'Vietnam', 'Hong',
       'Ireland', 'Hungary', 'Holand-Netherlands'], dtype=object)
```

```
df['native-country'].nunique()
```

```
42
```

How many people are working as tech support and have an annual income greater than 50k ?

```
df[df[' occupation']==' Tech-support']
```

	age	workclass	fnlwgt	education	education-num	\
24	59	Private	109015	HS-grad	9	
25	56	Local-gov	216851	Bachelors	13	
42	24	Private	172987	Bachelors	13	
55	43	Private	237993	Some-college	10	
64	29	Private	105598	Some-college	10	
...	
32396	56	Private	135458	HS-grad	9	
32457	33	Private	139057	Masters	14	
32546	37	Private	198216	Assoc-acdm	12	
32553	32	Private	116138	Masters	14	
32556	27	Private	257302	Assoc-acdm	12	

	marital-status	occupation	relationship	\
24	Divorced	Tech-support	Unmarried	
25	Married-civ-spouse	Tech-support	Husband	
42	Married-civ-spouse	Tech-support	Husband	
55	Married-civ-spouse	Tech-support	Husband	
64	Divorced	Tech-support	Not-in-family	
...	
32396	Divorced	Tech-support	Not-in-family	
32457	Married-civ-spouse	Tech-support	Husband	
32546	Divorced	Tech-support	Not-in-family	
32553	Never-married	Tech-support	Not-in-family	
32556	Married-civ-spouse	Tech-support	Wife	

	race	sex	capital-gain	capital-loss	\
24	White	Female	0	0	
25	White	Male	0	0	
42	White	Male	0	0	
55	White	Male	0	0	
64	White	Male	0	0	
...	
32396	Black	Female	0	0	
32457	Asian-Pac-Islander	Male	0	0	
32546	White	Female	0	0	
32553	Asian-Pac-Islander	Male	0	0	
32556	White	Female	0	0	

	hours-per-week	native-country	Annuual-Income
24	40	United-States	<=50K
25	40	United-States	>50K
42	50	United-States	<=50K
55	40	United-States	>50K
64	58	United-States	<=50K
...
32396	40	United-States	<=50K

32457	50	United-States	>50K
32546	40	United-States	<=50K
32553	11	Taiwan	<=50K
32556	38	United-States	<=50K

[927 rows x 15 columns]


```
df[df[' occupation']==' Tech-support'].count()
```

```
age          927
workclass    927
fnlwgt       927
education    927
education-num 927
marital-status 927
occupation   927
relationship 927
race         927
sex          927
capital-gain  927
capital-loss  927
hours-per-week 927
native-country 927
Annuual-Income 927
dtype: int64
```

#No of Males and females in the data

```
sns.countplot(df[' sex'])
plt.show()
```

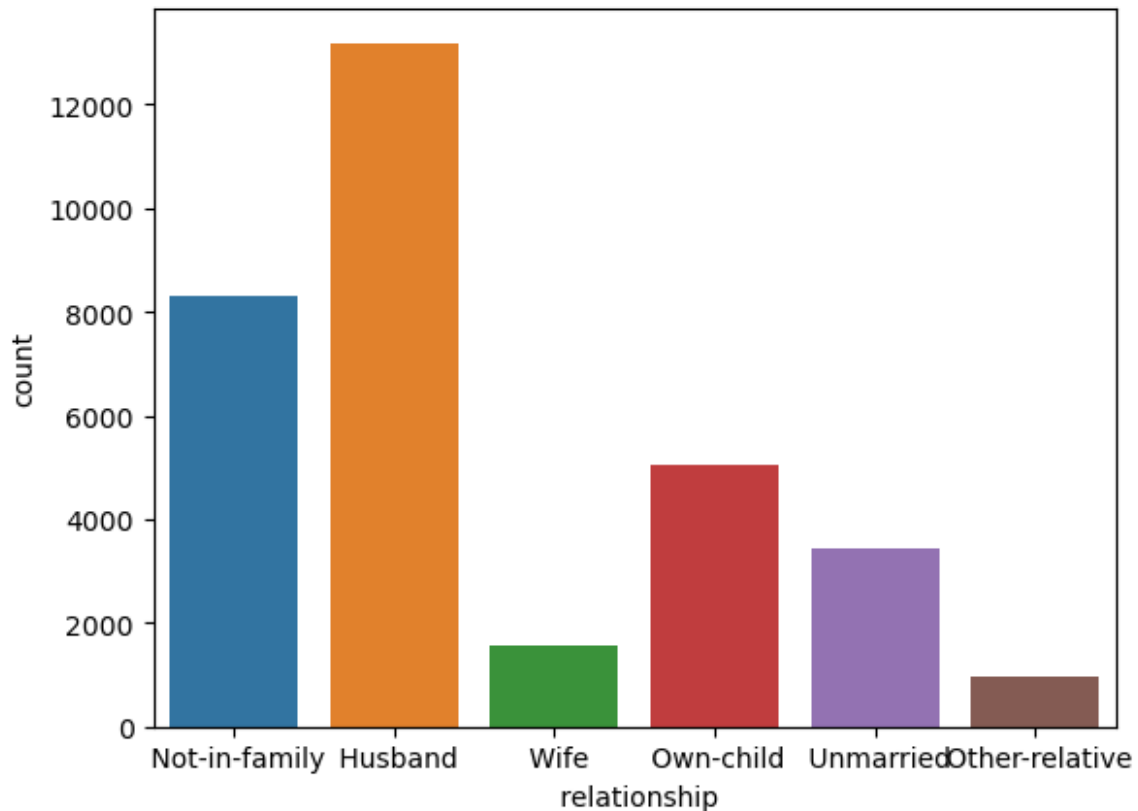
```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
_decorators.py:36: FutureWarning: Pass the following variable as a
keyword arg: x. From version 0.12, the only valid positional argument
will be `data`, and passing other arguments without an explicit
keyword will result in an error or misinterpretation.
  warnings.warn(
```



#How many people are married and un married in the census dataset

```
sns.countplot(df['relationship'])  
plt.show()
```


```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\  
_decorators.py:36: FutureWarning: Pass the following variable as a  
keyword arg: x. From version 0.12, the only valid positional argument  
will be `data`, and passing other arguments without an explicit  
keyword will result in an error or misinterpretation.  
warnings.warn(
```

#How many people are working in different sectors

```
plt.figure(figsize=(20,8))  
sns.countplot(df['workclass'])  
plt.show()
```

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\  
_decorators.py:36: FutureWarning: Pass the following variable as a  
keyword arg: x. From version 0.12, the only valid positional argument  
will be `data`, and passing other arguments without an explicit  
keyword will result in an error or misinterpretation.  
warnings.warn(
```



#Types of Education

```
df[' education'].unique()
```

```
array([' Bachelors', ' HS-grad', ' 11th', ' Masters', ' 9th',  
      ' Some-college', ' Assoc-acdm', ' Assoc-voc', ' 7th-8th',  
      ' Doctorate', ' Prof-school', ' 5th-6th', ' 10th', ' 1st-4th',  
      ' Preschool', ' 12th'], dtype=object)
```

How many people are having private work classes and are not from the United States of America?

```
df[' workclass'].unique()
```

```
array([' State-gov', ' Self-emp-not-inc', ' Private', ' Federal-gov',  
      ' Local-gov', ' ?', ' Self-emp-inc', ' Without-pay',  
      ' Never-worked'], dtype=object)
```

```
df[(df[' workclass']==' Private')&(df[' native-country']!=' United-  
States')]
```

	age	workclass	fnlwgt	education	education-num \
4	28	Private	338409	Bachelors	13
6	49	Private	160187	9th	5
14	40	Private	121772	Assoc-voc	11
15	34	Private	245487	7th-8th	4
35	48	Private	242406	11th	7
...
32508	45	Private	155093	10th	6
32510	39	Private	107302	HS-grad	9
32533	54	Private	337992	Bachelors	13
32547	43	Private	260761	HS-grad	9
32553	32	Private	116138	Masters	14
	marital-status		occupation		relationship \
4	Married-civ-spouse		Prof-specialty		Wife
6	Married-spouse-absent		Other-service		Not-in-family

14	Married-civ-spouse	Craft-repair	Husband
15	Married-civ-spouse	Transport-moving	Husband
35	Never-married	Machine-op-inspct	Unmarried
...
32508	Divorced	Other-service	Not-in-family
32510	Married-civ-spouse	Prof-specialty	Husband
32533	Married-civ-spouse	Exec-managerial	Husband
32547	Married-civ-spouse	Machine-op-inspct	Husband
32553	Never-married	Tech-support	Not-in-family

	race	sex	capital-gain	capital-loss	\
4	Black	Female	0	0	
6	Black	Female	0	0	
14	Asian-Pac-Islander	Male	0	0	
15	Amer-Indian-Eskimo	Male	0	0	
35	White	Male	0	0	
...	
32508	Black	Female	0	0	
32510	White	Male	0	0	
32533	Asian-Pac-Islander	Male	0	0	
32547	White	Male	0	0	
32553	Asian-Pac-Islander	Male	0	0	

	hours-per-week	native-country	Annuual-Income
4	40	Cuba	<=50K
6	16	Jamaica	<=50K
14	40	?	>50K
15	45	Mexico	<=50K
35	40	Puerto-Rico	<=50K
...
32508	38	Dominican-Republic	<=50K
32510	45	?	>50K
32533	50	Japan	>50K
32547	40	Mexico	<=50K
32553	11	Taiwan	<=50K

[2554 rows x 15 columns]

df[(df[' workclass']==' Private')&(df[' native-country']!=' United-States')][' workclass'].count()

2554

How many people are either having Annual Income(last column) less than or equal to 50k or their working hours is greater than or equal to 40 hrs

```
df[(df['Annuual-Income']==' <=50K')&(df['hours-per-week']>40)]
```

	age	workclass	fnlwgt	education	education-num
\					
13	32	Private	205019	Assoc-acdm	12

15	34	Private	245487	7th-8th	4
18	38	Private	28887	11th	7
28	39	Private	367260	HS-grad	9
30	23	Local-gov	190709	Assoc-acdm	12
...
32537	30	Private	345898	HS-grad	9
32543	45	Local-gov	119199	Assoc-acdm	12
32548	65	Self-emp-not-inc	99359	Prof-school	15
32550	43	Self-emp-not-inc	27242	Some-college	10
32552	43	Private	84661	Assoc-voc	11
	marital-status		occupation	relationship	\
13	Never-married		Sales	Not-in-family	
15	Married-civ-spouse		Transport-moving	Husband	
18	Married-civ-spouse		Sales	Husband	
28	Divorced		Exec-managerial	Not-in-family	
30	Never-married		Protective-serv	Not-in-family	
...	
32537	Never-married		Craft-repair	Not-in-family	
32543	Divorced		Prof-specialty	Unmarried	
32548	Never-married		Prof-specialty	Not-in-family	
32550	Married-civ-spouse		Craft-repair	Husband	
32552	Married-civ-spouse		Sales	Husband	
	race	sex	capital-gain	capital-loss	\
13	Black	Male	0	0	
15	Amer-Indian-Eskimo	Male	0	0	
18	White	Male	0	0	
28	White	Male	0	0	
30	White	Male	0	0	
...	
32537	Black	Male	0	0	
32543	White	Female	0	0	
32548	White	Male	1086	0	
32550	White	Male	0	0	
32552	White	Male	0	0	
	hours-per-week	native-country	Anuual-Income		
13	50	United-States	<=50K		
15	45	Mexico	<=50K		

18	50	United-States	<=50K
28	80	United-States	<=50K
30	52	United-States	<=50K
...
32537	46	United-States	<=50K
32543	48	United-States	<=50K
32548	60	United-States	<=50K
32550	50	United-States	<=50K
32552	45	United-States	<=50K

[5721 rows x 15 columns]

```
df[(df['Annuual-Income']== '<=50K') & (df[' hours-per-week']>40)].count()
```

```
age          5721
workclass    5721
fnlwgt       5721
education    5721
education-num 5721
marital-status 5721
occupation   5721
relationship 5721
race         5721
sex          5721
capital-gain 5721
capital-loss 5721
hours-per-week 5721
native-country 5721
Annuual-Income 5721
dtype: int64
```

Popoulation According to the country

```
l=df[' native-country'].unique()
print(l)
```

```
[' United-States' ' Cuba' ' Jamaica' ' India' ' ?' ' Mexico' ' South'
 ' Puerto-Rico' ' Honduras' ' England' ' Canada' ' Germany' ' Iran'
 ' Philippines' ' Italy' ' Poland' ' Columbia' ' Cambodia' ' Thailand'
 ' Ecuador' ' Laos' ' Taiwan' ' Haiti' ' Portugal' ' Dominican-
 Republic'
 ' El-Salvador' ' France' ' Guatemala' ' China' ' Japan' ' Yugoslavia'
 ' Peru' ' Outlying-US(Guam-USVI-etc)' ' Scotland' ' Trinidad&Tobago'
 ' Greece' ' Nicaragua' ' Vietnam' ' Hong' ' Ireland' ' Hungary'
 ' Holand-Netherlands']
```

```
d={}
for i in l:
    pop=df[df[' native-country']==i][' native-country'].count()
    d[i]=pop
population=pd.DataFrame(list(d.items()),columns=['Native
```

```
country', 'Population'])
population.head(60)
```

	Native country	Population
0	United-States	29153
1	Cuba	95
2	Jamaica	81
3	India	100
4	?	582
5	Mexico	639
6	South	80
7	Puerto-Rico	114
8	Honduras	13
9	England	90
10	Canada	121
11	Germany	137
12	Iran	43
13	Philippines	198
14	Italy	73
15	Poland	60
16	Columbia	59
17	Cambodia	19
18	Thailand	18
19	Ecuador	28
20	Laos	18
21	Taiwan	51
22	Haiti	44
23	Portugal	37
24	Dominican-Republic	70
25	El-Salvador	106
26	France	29
27	Guatemala	62
28	China	75
29	Japan	62
30	Yugoslavia	16
31	Peru	31
32	Outlying-US(Guam-USVI-etc)	14
33	Scotland	12
34	Trinidad&Tobago	19
35	Greece	29
36	Nicaragua	34
37	Vietnam	67
38	Hong	20
39	Ireland	24
40	Hungary	13
41	Holand-Netherlands	1

Observations from table

- United-States has Highest Population

- Holand-Netherlands has lowest Population

```
df['Annuual-Income'].unique()
array([' <=50K', ' >50K'], dtype=object)

from sklearn.preprocessing import LabelEncoder
le=LabelEncoder

from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()

# Fit and transform the 'workclass' column
df[' workclass'] = le.fit_transform(df[' workclass'])
df[' fnlwgt'] = le.fit_transform(df[' fnlwgt'])
df[' education'] = le.fit_transform(df[' education'])
df[' marital-status'] = le.fit_transform(df[' marital-status'])
df[' occupation'] = le.fit_transform(df[' occupation'])
df[' relationship'] = le.fit_transform(df[' relationship'])
df[' sex'] = le.fit_transform(df[' sex'])
df[' race'] = le.fit_transform(df[' race'])
df[' native-country'] = le.fit_transform(df[' native-country'])
df['Annuual-Income'] = le.fit_transform(df['Annuual-Income'])

df.head()
```

	age	workclass	fnlwgt	education	education-num	marital-
status \						
0	39	7	2671	9	13	
4						
1	50	6	2926	9	13	
2						
2	38	4	14086	11	9	
0						
3	53	4	15336	1	7	
2						
4	28	4	19355	9	13	
2						

	occupation	relationship	race	sex	capital-gain	capital-
loss \						
0	1	1	4	1	2174	
0						
1	4	0	4	1	0	
0						
2	6	1	4	1	0	
0						
3	6	0	2	1	0	
0						
4	10	5	2	0	0	

0

	hours-per-week	native-country	Annuual-Income
0	40	39	0
1	13	39	0
2	40	39	0
3	40	39	0
4	40	5	0

```
df.drop(['age'],axis=1,inplace=True)
```

df

	workclass	fnlwgt	education	education-num	marital-
status \					
0	7	2671	9	13	
4					
1	6	2926	9	13	
2					
2	4	14086	11	9	
0					
3	4	15336	1	7	
2					
4	4	19355	9	13	
2					
...
.					
32556	4	16528	7	12	
2					
32557	4	8080	11	9	
2					
32558	4	7883	11	9	
6					
32559	4	12881	11	9	
4					
32560	5	17825	11	9	
2					
	occupation	relationship	race	sex	capital-gain
capital-loss \					
0	1	1	4	1	2174
0					
1	4	0	4	1	0
0					
2	6	1	4	1	0
0					
3	6	0	2	1	0
0					
4	10	5	2	0	0
0					


```

...
...
32556      13      5      4      0      0
0
32557      7      0      4      1      0
0
32558      1      4      4      0      0
0
32559      1      3      4      1      0
0
32560      4      5      4      0     15024
0

      hours-per-week  native-country  Annual-Income
0              40              39              0
1              13              39              0
2              40              39              0
3              40              39              0
4              40              5              0
...
32556      38              39              0
32557      40              39              1
32558      40              39              0
32559      20              39              0
32560      40              39              1

[32537 rows x 14 columns]

```

Segregation of Dependent and Independent Columns

```

x=df.drop('Annual-Income',axis=1)
y=df['Annual-Income']

```

```

x.head(20)

```

```

      workclass  fnlwgt  education  education-num  marital-
status \
0          7    2671         9         13         4
1          6    2926         9         13         2
2          4   14086        11          9         0
3          4   15336         1          7         2
4          4   19355         9         13         2

```

5	4	17700	12	14	2
6	4	8536	6	5	3
7	6	13620	11	9	2
8	4	1318	12	14	4
9	4	8460	9	13	2
10	4	17530	15	10	2
11	7	7077	9	13	2
12	4	5772	9	13	4
13	4	13217	7	12	4
14	4	5722	8	11	2
15	4	15963	5	4	2
16	6	10163	11	9	4
17	4	11257	11	9	4
18	4	314	1	7	2
19	6	17974	12	14	0
loss \ occupation relationship race sex capital-gain capital-					
0	1	1	4	1	2174
0					
1	4	0	4	1	0
0					
2	6	1	4	1	0
0					
3	6	0	2	1	0
0					
4	10	5	2	0	0
0					
5	4	5	4	0	0
0					
6	8	1	2	0	0
0					
7	4	0	4	1	0
0					
8	10	1	4	0	14084
0					

9	4	0	4	1	5178
0					
10	4	0	2	1	0
0					
11	10	0	1	1	0
0					
12	1	3	4	0	0
0					
13	12	1	2	1	0
0					
14	3	0	1	1	0
0					
15	14	0	0	1	0
0					
16	5	3	4	1	0
0					
17	7	4	4	1	0
0					
18	12	0	4	1	0
0					
19	4	4	4	0	0
0					

	hours-per-week	native-country
0	40	39
1	13	39
2	40	39
3	40	39
4	40	5
5	40	39
6	16	23
7	45	39
8	50	39
9	40	39
10	80	39
11	40	19
12	30	39
13	50	39
14	40	0
15	45	26
16	35	39
17	40	39
18	50	39
19	45	39

```
y.head(20)
```

0	0
1	0
2	0

```

3      0
4      0
5      0
6      0
7      1
8      1
9      1
10     1
11     1
12     0
13     0
14     1
15     0
16     0
17     0
18     0
19     1
Name: Annual-Income, dtype: int32

```

Logistic Regression Model

```

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.4, random_state=42)
from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
model.fit(x_train,y_train)
pred=model.predict(x_test)
from sklearn.linear_model import
LogisticRegression
model=LogisticRegression()
model.fit(x_train,y_train)
pred=model.predict(x_test)

```

C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
from sklearn.metrics import *
print("Accuracy",accuracy_score(y_test,pred))
print("recall",recall_score(y_test,pred))
```

```
Accuracy 0.8086822896657703
recall 0.3857052896725441
```


```
sns.distplot(y_test,hist=False,color="red")
sns.distplot(pred,hist=False,color="blue")
plt.show()
```

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
distributions.py:2619: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar
flexibility) or `kdeplot` (an axes-level function for kernel density
plots).
```

```
warnings.warn(msg, FutureWarning)
```

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
distributions.py:2619: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar
flexibility) or `kdeplot` (an axes-level function for kernel density
plots).
```

```
warnings.warn(msg, FutureWarning)
```



_____Hyperparameter Tunning of Logistic Regression _____

```
logistic_hyperparameter = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2'],
    'solver': ['liblinear', 'saga']
}

logistic_hyper= RandomizedSearchCV(
    estimator=LogisticRegression(random_state=31),
    param_distributions=logistic_hyperparameter,
    n_jobs=-1,
    n_iter=100,
    cv=3,
    verbose=3,
    scoring='accuracy',
    random_state=0
)

logistic_hyper.fit(x_train, y_train)
```

```

print("Logistic Regression Best Score:", logistic_hyper.best_score_)
print("Logistic Regression Best Params:", logistic_hyper.best_params_)

C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_search.py:307: UserWarning: The total space of
parameters 24 is smaller than n_iter=100. Running 24 iterations. For
exhaustive searches, use GridSearchCV.
  warnings.warn(

Fitting 3 folds for each of 24 candidates, totalling 72 fits
Logistic Regression Best Score: 0.8249155079142386
Logistic Regression Best Params: {'solver': 'liblinear', 'penalty':
'l1', 'C': 1}

model=LogisticRegression(solver='liblinear',penalty='l1',C=1)
model.fit(x_train,y_train)
pred=model.predict(x_test)
print("Accuracy",accuracy_score(y_test,pred))
print("recall",recall_score(y_test,pred))

Accuracy 0.8237418363426815
recall 0.4474181360201511

```

Decision Tree Model

```

from sklearn.tree import DecisionTreeClassifier
model=DecisionTreeClassifier()
model.fit(x_train,y_train)
pred=model.predict(x_test)

print("Accuracy",accuracy_score(y_test,pred))
print("recall",recall_score(y_test,pred))

Accuracy 0.8026892047637342
recall 0.6026448362720404


sns.distplot(y_test,hist=False,color="red")
sns.distplot(pred,hist=False,color="blue")
plt.show()

C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
distributions.py:2619: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar
flexibility) or `kdeplot` (an axes-level function for kernel density
plots).
  warnings.warn(msg, FutureWarning)
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
distributions.py:2619: FutureWarning: `distplot` is a deprecated

```

function and will be removed in a future version. Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``kdeplot`` (an axes-level function for kernel density plots).

```
warnings.warn(msg, FutureWarning)
```



Hyperparameter Tuning of Decision Tree Classifier

```
from sklearn.model_selection import RandomizedSearchCV
decisionhyperparameter={
    'max_depth': [None, range(1, 20)],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt', 'log2'],
    'criterion': ['gini', 'entropy']
}
decision_hyper=RandomizedSearchCV(estimator=DecisionTreeClassifier(random_state=31),
```



```

param_distributions=decisionhyperparameter,
                    n_jobs=-1,
                    n_iter=100,
                    cv=3,
                    verbose=3,
                    scoring='accuracy',
                    random_state=0
                )
decision_hyper.fit(x_train,y_train)
print(decision_hyper.best_score_)
print(decision_hyper.best_params_)

```

Fitting 3 folds for each of 100 candidates, totalling 300 fits
0.8438170001278626
{'min_samples_split': 10, 'min_samples_leaf': 4, 'max_features':
'log2', 'max_depth': None, 'criterion': 'gini'}

C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection_validation.py:425: FitFailedWarning:
204 fits failed out of a total of 300.
The score on these train-test partitions for these parameters will be
set to nan.
If these failures are not expected, you can try to debug them by
setting error_score='raise'.

Below are more details about the failures:

```

-----
-----
150 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_validation.py", line 729, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 1145, in wrapper
    estimator._validate_params()
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 638, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_depth'
parameter of DecisionTreeClassifier must be an int in the range [1,
inf) or None. Got range(1, 20) instead.

```

```

-----
-----
34 fits failed with the following error:

```

```

Traceback (most recent call last):
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_validation.py", line 729, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 1145, in wrapper
    estimator._validate_params()
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 638, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of DecisionTreeClassifier must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'sqrt', 'log2'} or None. Got 'auto' instead.

```


20 fits failed with the following error:

```

Traceback (most recent call last):
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_validation.py", line 729, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 1145, in wrapper
    estimator._validate_params()
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 638, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of DecisionTreeClassifier must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'log2', 'sqrt'} or None. Got 'auto' instead.

```

```

    warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_search.py:979: UserWarning: One or more of the test
scores are non-finite: [          nan  0.81917842  0.83920705          nan
0.84156318          nan
          nan          nan  0.84156318          nan  0.83597968  0.82952564
          nan          nan  0.83101104          nan  0.82952564  0.81733414
          nan          nan          nan          nan  0.843817          nan

```

```

nan nan 0.83920705 nan nan nan
0.84356099 nan nan nan nan nan
nan 0.81016295 nan nan 0.81088001 0.84294637
nan nan 0.84356099 nan nan nan
0.84156318 0.843817 nan nan nan
nan nan nan nan nan
nan 0.82998695 nan nan 0.83992418 nan
nan nan nan nan nan
0.83992418 0.82665717 0.81917842 nan nan nan
nan 0.84356099 0.84294637 0.82998695 nan 0.81088001
0.83597968 0.84156318 nan nan nan
nan nan 0.82665717 nan 0.83413573 nan
nan 0.84356099 nan nan]
warnings.warn(

```

```

model=DecisionTreeClassifier(min_samples_split= 10, min_samples_leaf=
4, max_features= 'log2', max_depth= None, criterion= 'gini')
model.fit(x_train,y_train)
pred=model.predict(x_test)
print("Accuracy",accuracy_score(y_test,pred))
print("recall",recall_score(y_test,pred))

Accuracy 0.8341913177103343
recall 0.5639168765743073

```

Random Forest Model

```

from sklearn.ensemble import RandomForestClassifier
model=RandomForestClassifier()
model.fit(x_train,y_train)
pred=model.predict(x_test)

```

```

print("Accuracy",accuracy_score(y_test,pred))
print("recall",recall_score(y_test,pred))

```

```

Accuracy 0.8375720322704572
recall 0.6001259445843828

```

```

sns.distplot(y_test,hist=False,color="red")
sns.distplot(pred,hist=False,color="blue")
plt.show()


```

```

C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
distributions.py:2619: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar
flexibility) or `kdeplot` (an axes-level function for kernel density
plots).
  warnings.warn(msg, FutureWarning)

```

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
distributions.py:2619: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar
flexibility) or `kdeplot` (an axes-level function for kernel density
plots).
  warnings.warn(msg, FutureWarning)
```



_____Hyperparameter Tunning of Random Forest Classifier_____

```
from sklearn.model_selection import RandomizedSearchCV
Randomforesthyperparameter={
    'n_estimators':[50,100,200,300,400,500],
    'max_depth':[None,range(1,20)],
    'min_samples_split':[2,5,10],
    'min_samples_leaf':[1,2,4],
    'max_features':['auto','sqrt','log2']
}
```

```

Random_hyper=RandomizedSearchCV(estimator=RandomForestClassifier(rando
m_state=31),

param_distributions=Randomforesthyperparameter,
                                n_jobs=-1,
                                n_iter=100,
                                cv=3,
                                verbose=3,
                                scoring='accuracy',
                                random_state=0
                                )
Random_hyper.fit(x_train,y_train)
print(Random_hyper.best_score_)
print(Random_hyper.best_params_)

```

Fitting 3 folds for each of 100 candidates, totalling 300 fits

```

C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_validation.py:425: FitFailedWarning:
192 fits failed out of a total of 300.
The score on these train-test partitions for these parameters will be
set to nan.
If these failures are not expected, you can try to debug them by
setting error_score='raise'.

```

Below are more details about the failures:

```

-----
-----
138 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_validation.py", line 729, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 1145, in wrapper
    estimator._validate_params()
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 638, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_depth'
parameter of RandomForestClassifier must be an int in the range [1,
inf) or None. Got range(1, 20) instead.

```

```

-----
-----
40 fits failed with the following error:

```

```

Traceback (most recent call last):
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_validation.py", line 729, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 1145, in wrapper
    estimator._validate_params()
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 638, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of RandomForestClassifier must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'log2', 'sqrt'} or None. Got 'auto' instead.

```


14 fits failed with the following error:

```

Traceback (most recent call last):
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_validation.py", line 729, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 1145, in wrapper
    estimator._validate_params()
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 638, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of RandomForestClassifier must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'sqrt', 'log2'} or None. Got 'auto' instead.

```

```

    warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_search.py:979: UserWarning: One or more of the test
scores are non-finite: [          nan 0.85477914          nan          nan
0.86067002 0.85908201
0.85923566          nan 0.85421566 0.85949179          nan          nan
0.85903079          nan          nan 0.85488156          nan          nan
0.85892832 0.84714657          nan 0.85687933 0.85887709          nan

```

```

0.86010657      nan      nan      nan      nan      nan
      nan      nan      nan      nan      nan      nan
0.84724912      nan 0.85682817 0.85339608 0.85841608      nan
0.85856984 0.85324241 0.85979927 0.86036274 0.8593895      nan
      nan      nan      nan      nan 0.84786382 0.85682817
      nan 0.86010657      nan 0.84765895      nan      nan
      nan      nan 0.858826      nan      nan      nan
      nan      nan      nan      nan      nan      nan
      nan      nan 0.85892832 0.85856984 0.85938935      nan
0.85687933      nan 0.85918446      nan      nan      nan
0.85928688      nan      nan 0.86000413 0.85713553      nan
      nan      nan 0.85687938      nan      nan      nan
      nan      nan      nan      nan]
warnings.warn(
0.8606700152109527
{'n_estimators': 200, 'min_samples_split': 10, 'min_samples_leaf': 2,
'max_features': 'log2', 'max_depth': None}

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, recall_score

# Assuming x_train, y_train, x_test, and y_test are defined

model = RandomForestClassifier(n_estimators=200, min_samples_split=10,
min_samples_leaf=2, max_features='log2', max_depth=None)
model.fit(x_train, y_train)
pred = model.predict(x_test)

print("Accuracy:", accuracy_score(y_test, pred))
print("Recall:", recall_score(y_test, pred))

Accuracy: 0.8559354590856704
Recall: 0.5954030226700252

```

Naive Bayes Model

```

from sklearn.naive_bayes import GaussianNB

model=GaussianNB()
model.fit(x_train,y_train)
pred=model.predict(x_test)

print("Accuracy",accuracy_score(y_test,pred))
print("recall",recall_score(y_test,pred))

Accuracy 0.7975412985017287
recall 0.3195843828715365

```


```
sns.distplot(y_test,hist=False,color="red")
sns.distplot(pred,hist=False,color="blue")
plt.show()
```

C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

```
warnings.warn(msg, FutureWarning)
```

C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

```
warnings.warn(msg, FutureWarning)
```



KNN Model

```
from sklearn.neighbors import KNeighborsClassifier
```



```
model=KNeighborsClassifier(n_neighbors=3)
model.fit(x_train,y_train)
prediction=model.predict(x_test)
```

```
print("Accuracy",accuracy_score(y_test,pred))
print("recall",recall_score(y_test,pred))
```

```
Accuracy 0.8237418363426815
recall 0.4474181360201511
```


```
sns.distplot(y_test,hist=False,color="red")
sns.distplot(pred,hist=False,color="blue")
plt.show()
```

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
distributions.py:2619: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar
flexibility) or `kdeplot` (an axes-level function for kernel density
plots).
```

```
warnings.warn(msg, FutureWarning)
```

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
distributions.py:2619: FutureWarning: `distplot` is a deprecated
function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar
flexibility) or `kdeplot` (an axes-level function for kernel density
plots).
```

```
warnings.warn(msg, FutureWarning)
```



Hyperparameter Tuning of KNN

```
from sklearn.model_selection import RandomizedSearchCV
knnhyperparameter={
    'n_neighbors': [3,5,7,9,11],
    'weights': ['uniform', 'distance'],
    'metric': ['euclidean', 'manhattan']
}
knn_hyper=RandomizedSearchCV(estimator=KNeighborsClassifier(),
                             param_distributions=knnhyperparameter,
                             n_jobs=-1,
                             n_iter=100,
                             cv=3,
                             verbose=3,
                             scoring='accuracy',
                             random_state=0
                             )
knn_hyper.fit(x_train,y_train)
print(knn_hyper.best_score_)
print(knn_hyper.best_params_)
```

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_search.py:307: UserWarning: The total space of
parameters 20 is smaller than n_iter=100. Running 20 iterations. For
exhaustive searches, use GridSearchCV.
```

```
warnings.warn(
```

```
Fitting 3 folds for each of 20 candidates, totalling 60 fits
```

```
0.8093433695835167
```

```
{'weights': 'distance', 'n_neighbors': 11, 'metric': 'manhattan'}
```

```
model=KNeighborsClassifier(weights= 'distance', n_neighbors= 11,
metric= 'manhattan')
```

```
model.fit(x_train,y_train)
```

```
prediction=model.predict(x_test)
```

```
print("Accuracy",accuracy_score(y_test,pred))
```

```
print("recall",recall_score(y_test,pred))
```

```
Accuracy 0.8237418363426815
```

```
recall 0.4474181360201511
```

- Among above Models Random forest gives best accuracy than other Models