

Importing Necessary Packages

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.model_selection import GridSearchCV
```

Data Preprocessing

```
df=pd.read_csv(r"C:\Users\Arigala.Adarsh\Downloads\diabetes(1).csv")
```

```
df.head(10)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	
BMI \						
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
5	5	116	74	0	0	25.6
6	3	78	50	32	88	31.0
7	10	115	0	0	0	35.3
8	2	197	70	45	543	30.5
9	8	125	96	0	0	0.0

DiabetesPedigreeFunction Age Outcome

0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
5	0.201	30	0
6	0.248	26	1
7	0.134	29	0
8	0.158	53	1
9	0.232	54	1

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

```
dtypes: float64(2), int64(7)
```

```
memory usage: 54.1 KB
```

```
df.columns
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
       'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

Exploratory Data Analysis

```
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness
Insulin \				
count	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458
std	3.369578	31.972618	19.355807	15.952218
115.244002				

min	0.000000	0.000000	0.000000	0.000000
0.000000				
25%	1.000000	99.000000	62.000000	0.000000
0.000000				
50%	3.000000	117.000000	72.000000	23.000000
30.500000				
75%	6.000000	140.250000	80.000000	32.000000
127.250000				
max	17.000000	199.000000	122.000000	99.000000
846.000000				

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
df.shape
```

```
(768, 9)
```

```
df.isnull().sum()
```

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0

```
dtype: int64
```


```
sns.countplot(df.Outcome)
```

```
plt.title('Count 0 & 1')
```

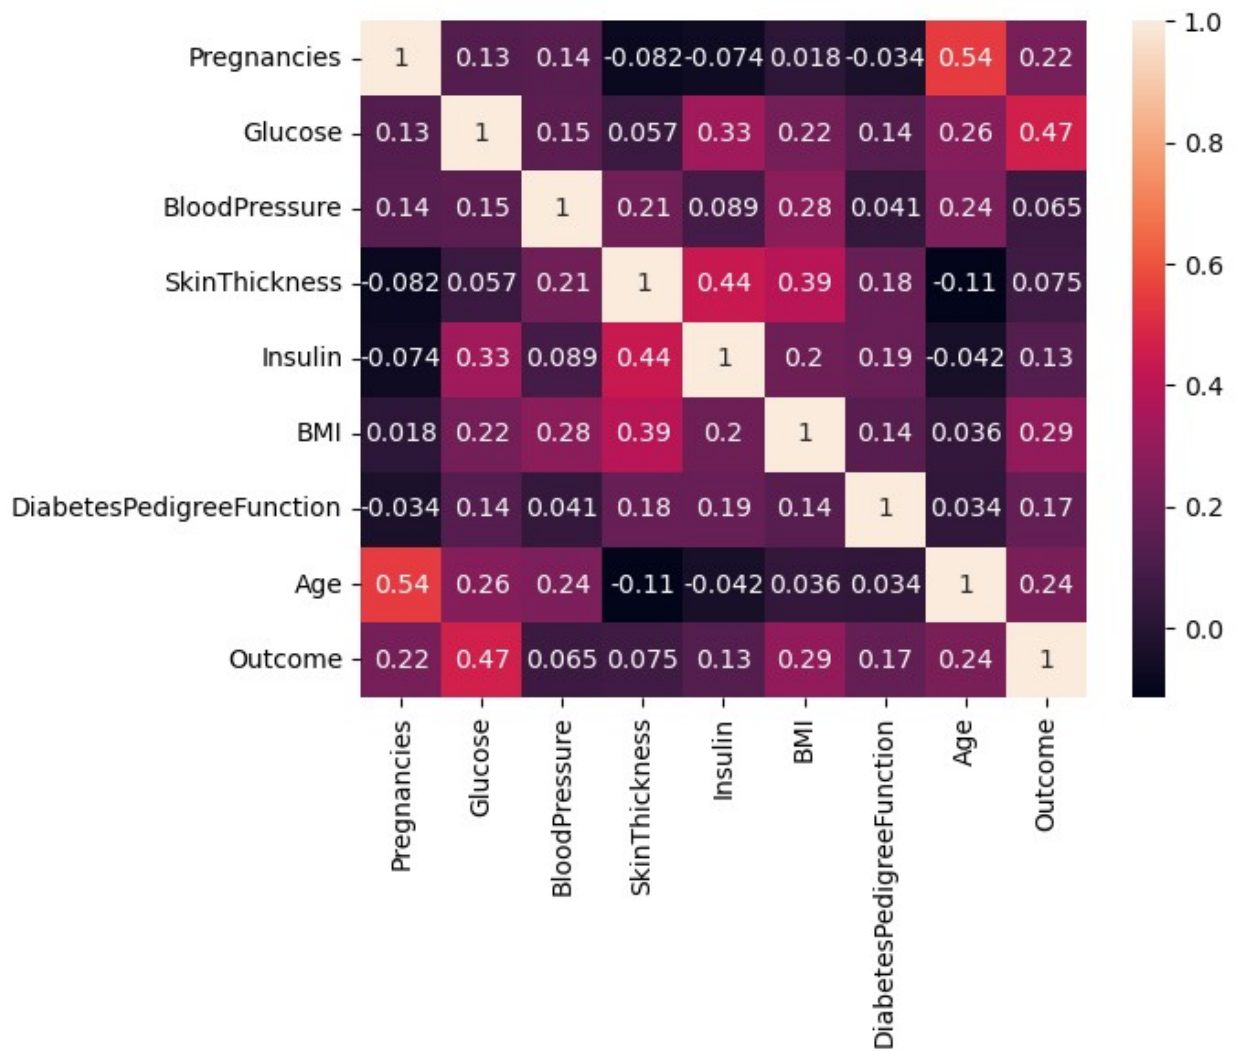
```
plt.show()
```

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
_decorators.py:36: FutureWarning: Pass the following variable as a
keyword arg: x. From version 0.12, the only valid positional argument
will be `data`, and passing other arguments without an explicit
keyword will result in an error or misinterpretation.
```

```
warnings.warn(
```

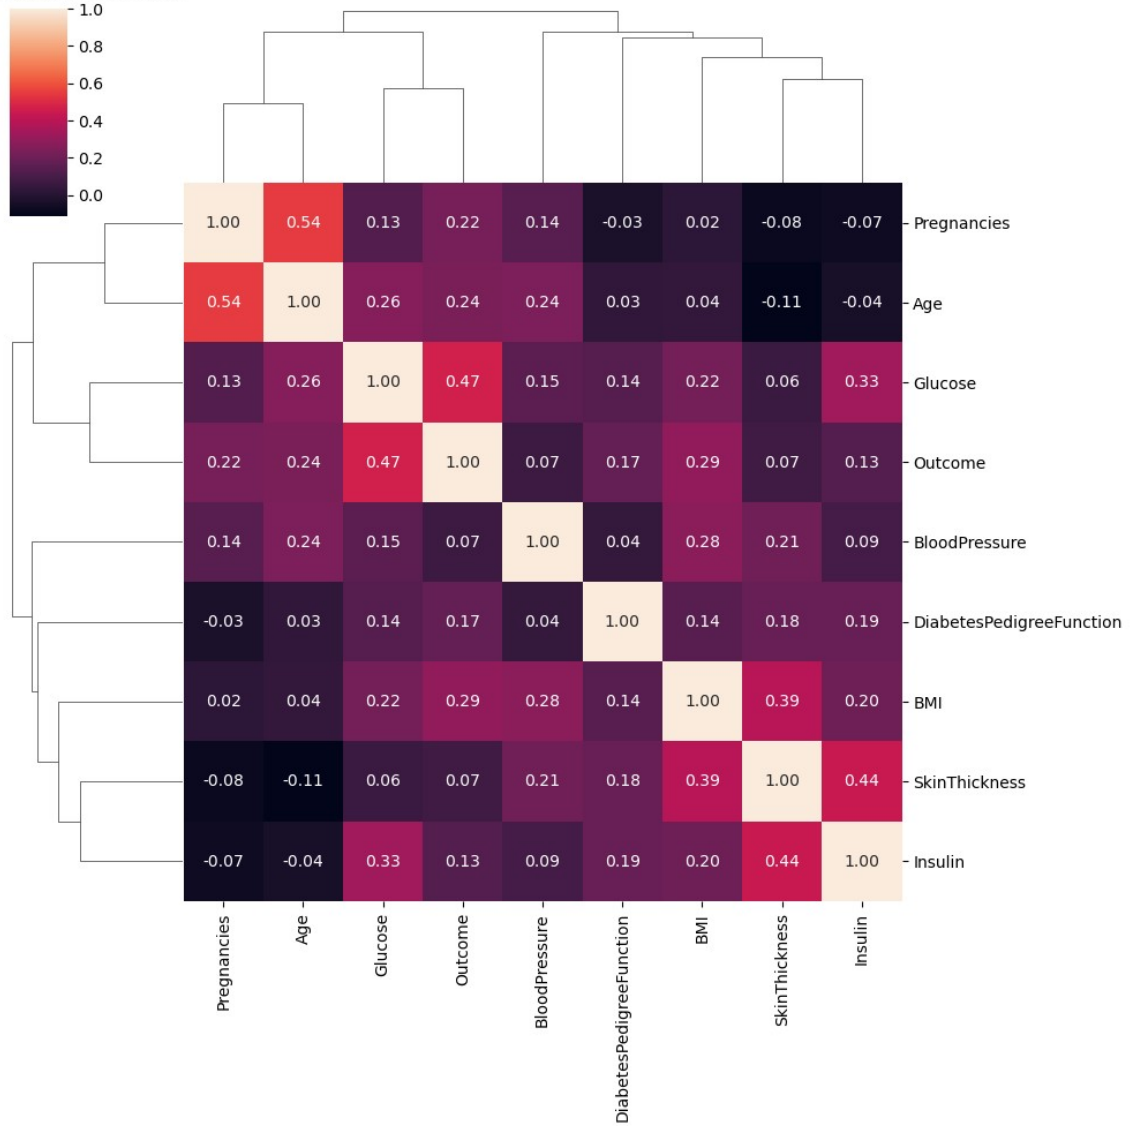


```
sns.heatmap(df.corr(),annot=True)  
plt.show()
```

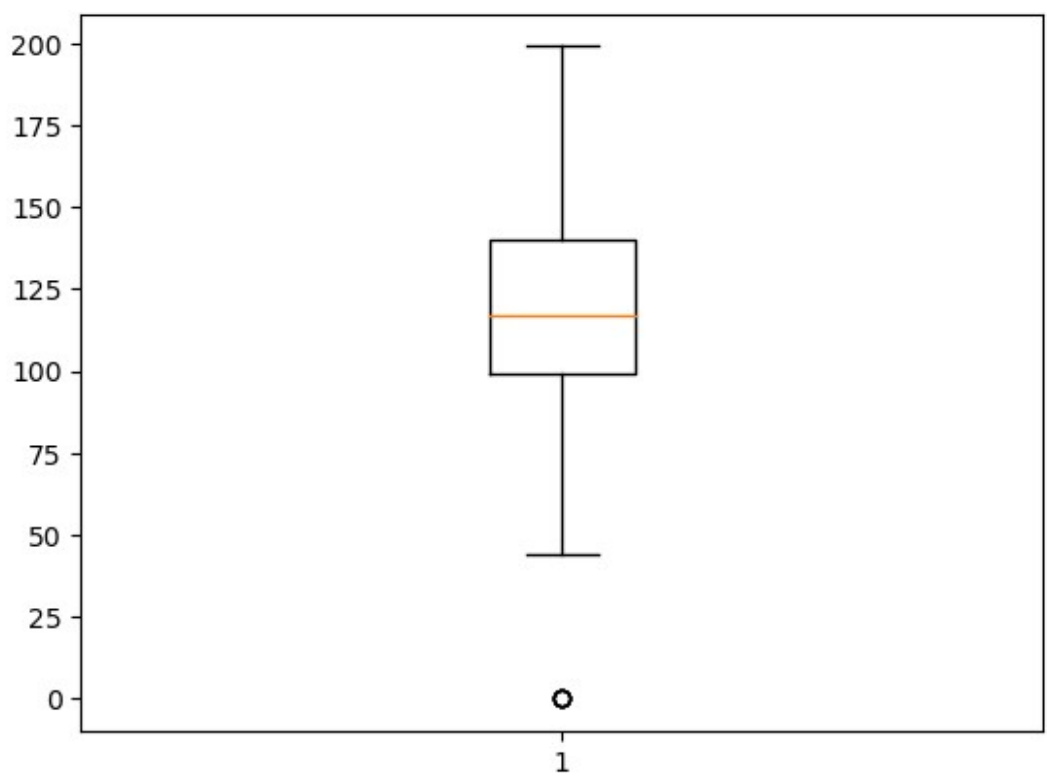
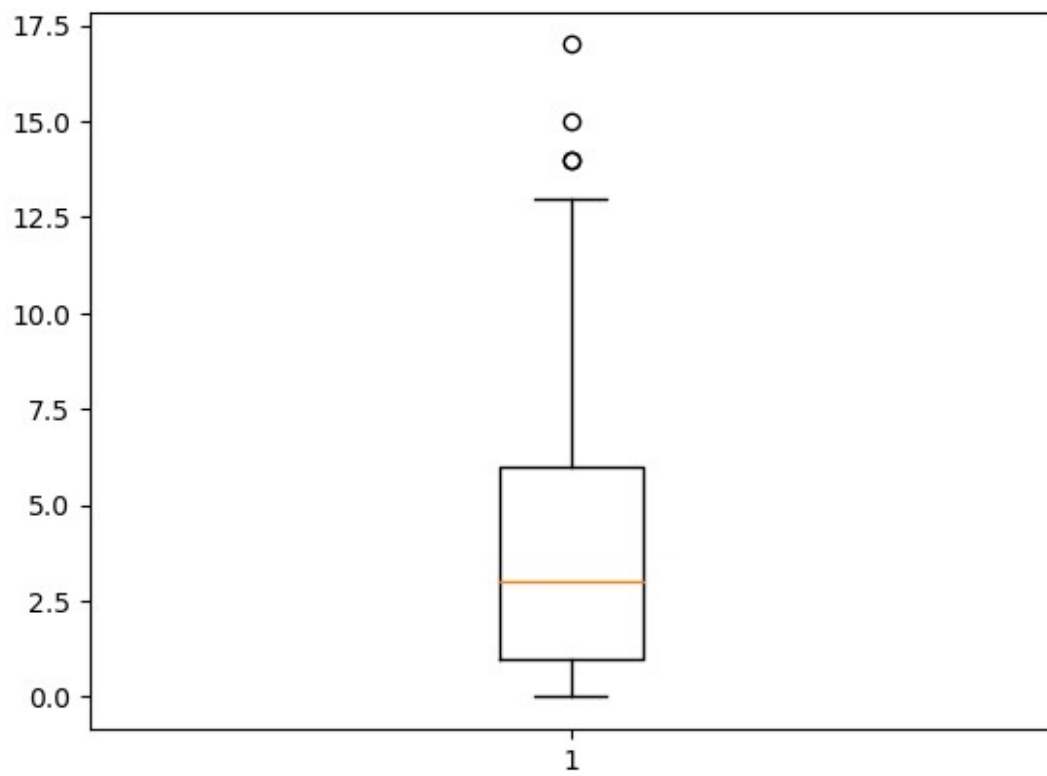



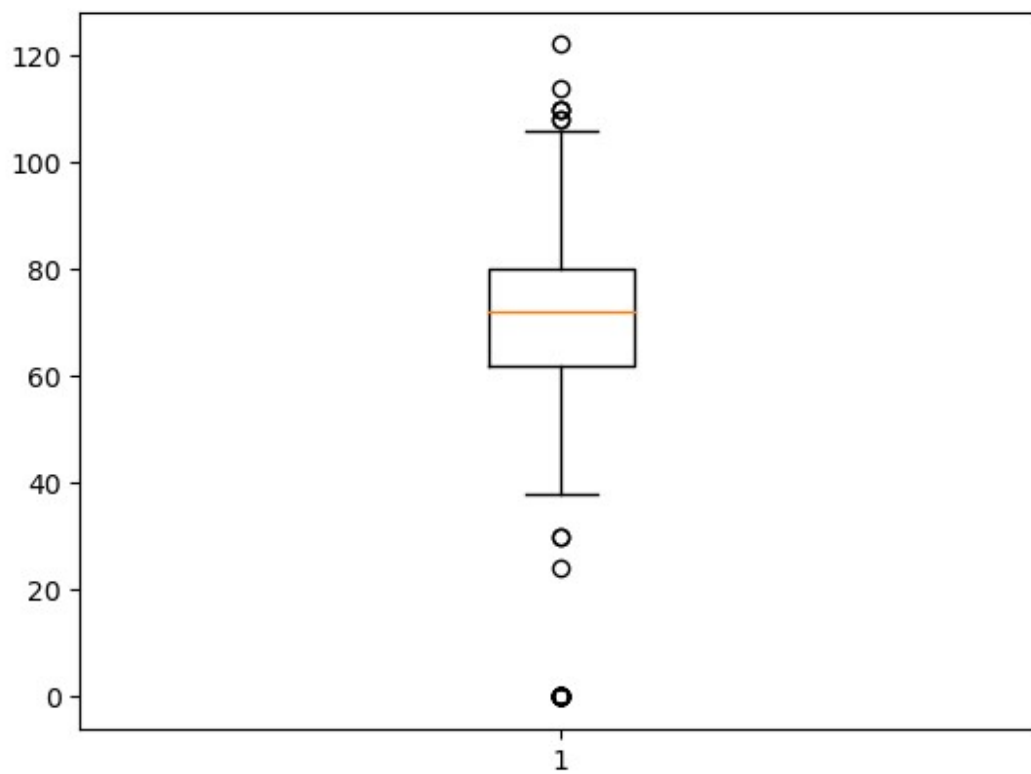
```
sns.clustermap(df.corr(),annot=True,fmt=".2f")
plt.title('Correlation Between Features')
plt.show()
```

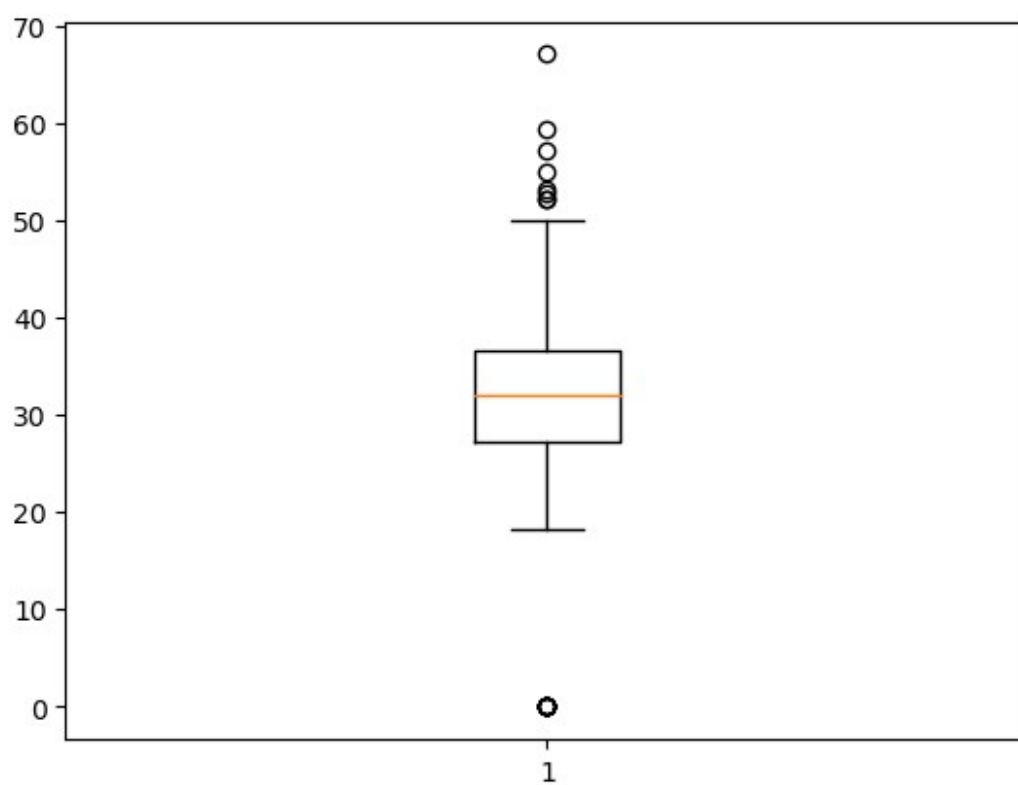
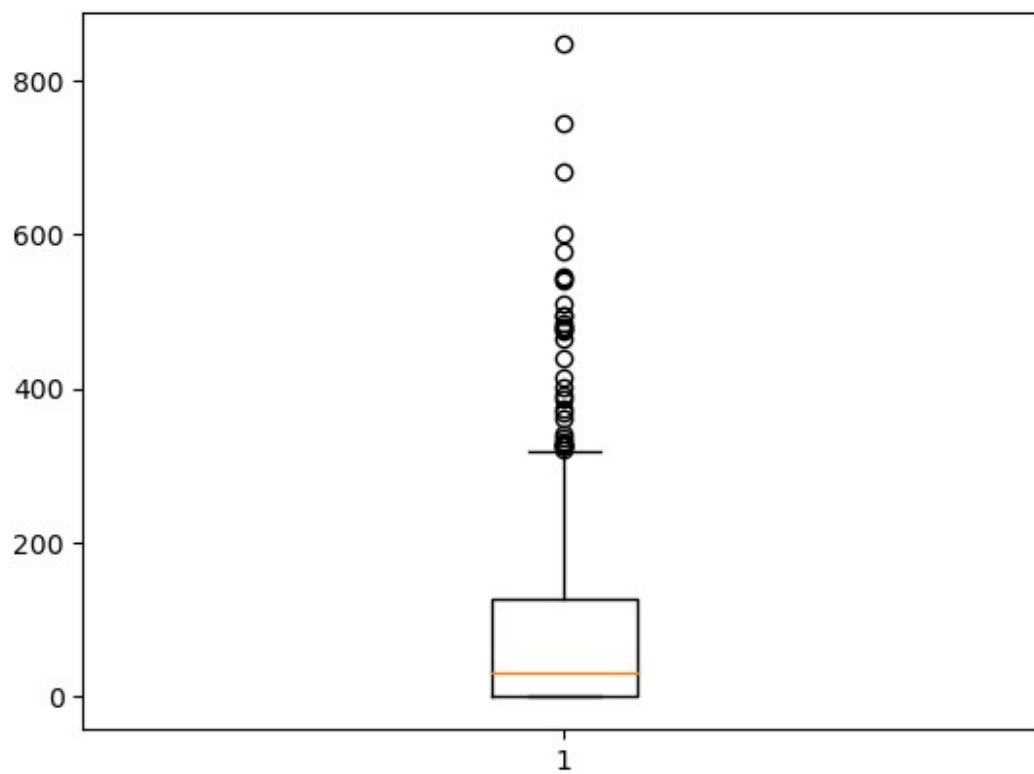
Correlation Between Features





```
for i in df.columns:
    if(i!='Outcome'):
        plt.boxplot(df[i])
        plt.show()
```









```

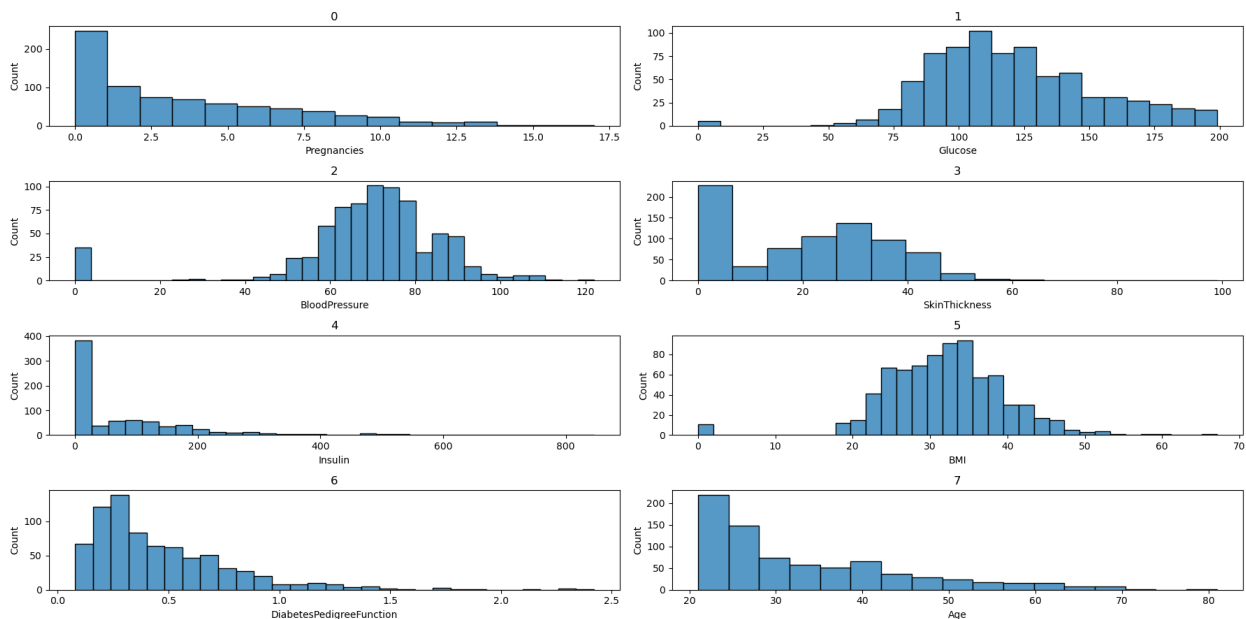
data1 = df[df["Outcome"] == 1]
columns = df.columns[:8]

plt.subplots(figsize=(18, 9))

for i, column in enumerate(columns):
    ax= plt.subplot(4, 2, i + 1)
    # Plot the histogram for all data
    sns.histplot(df[column], kde=False)
    plt.title(i)

plt.tight_layout()
plt.show()

```

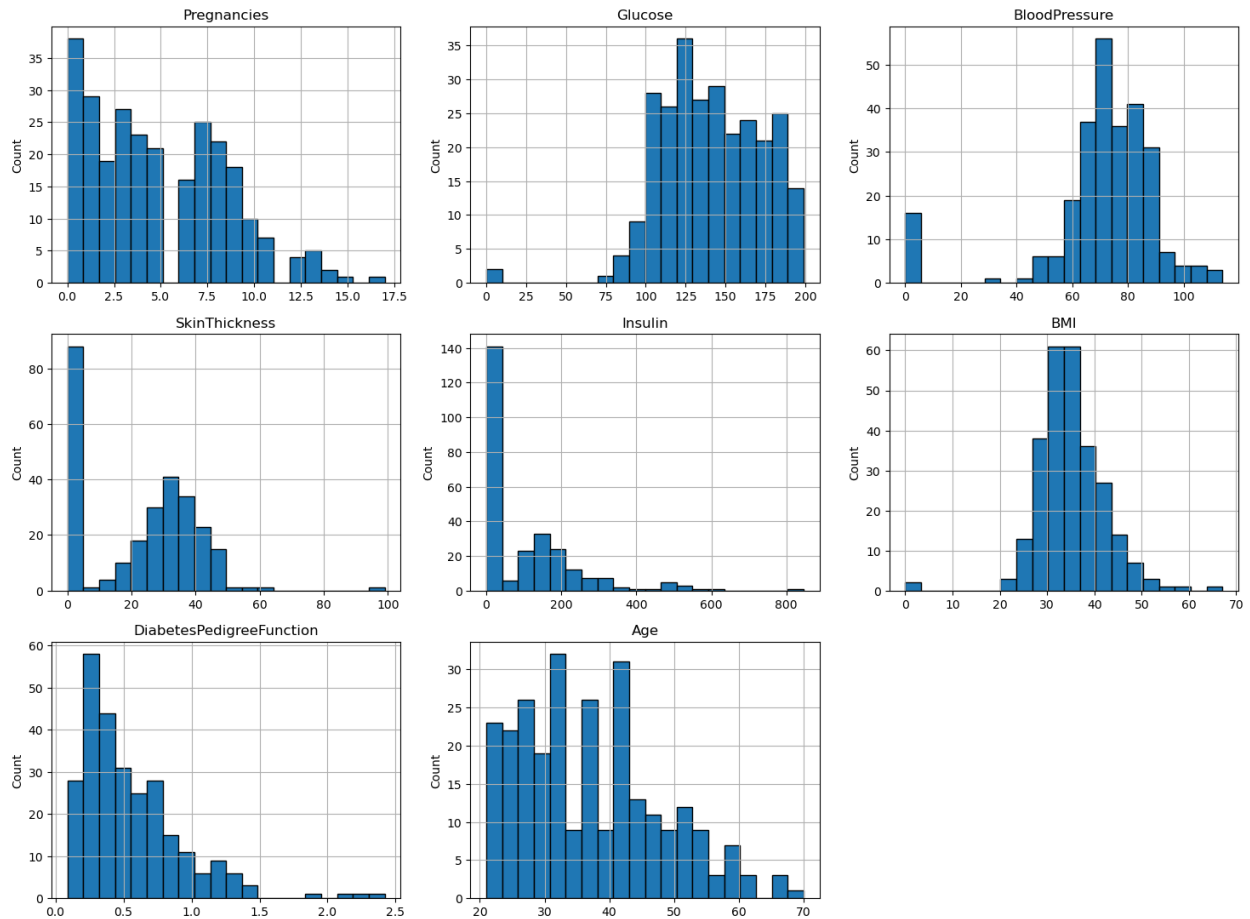


Analysis of Diabetic Cases

```

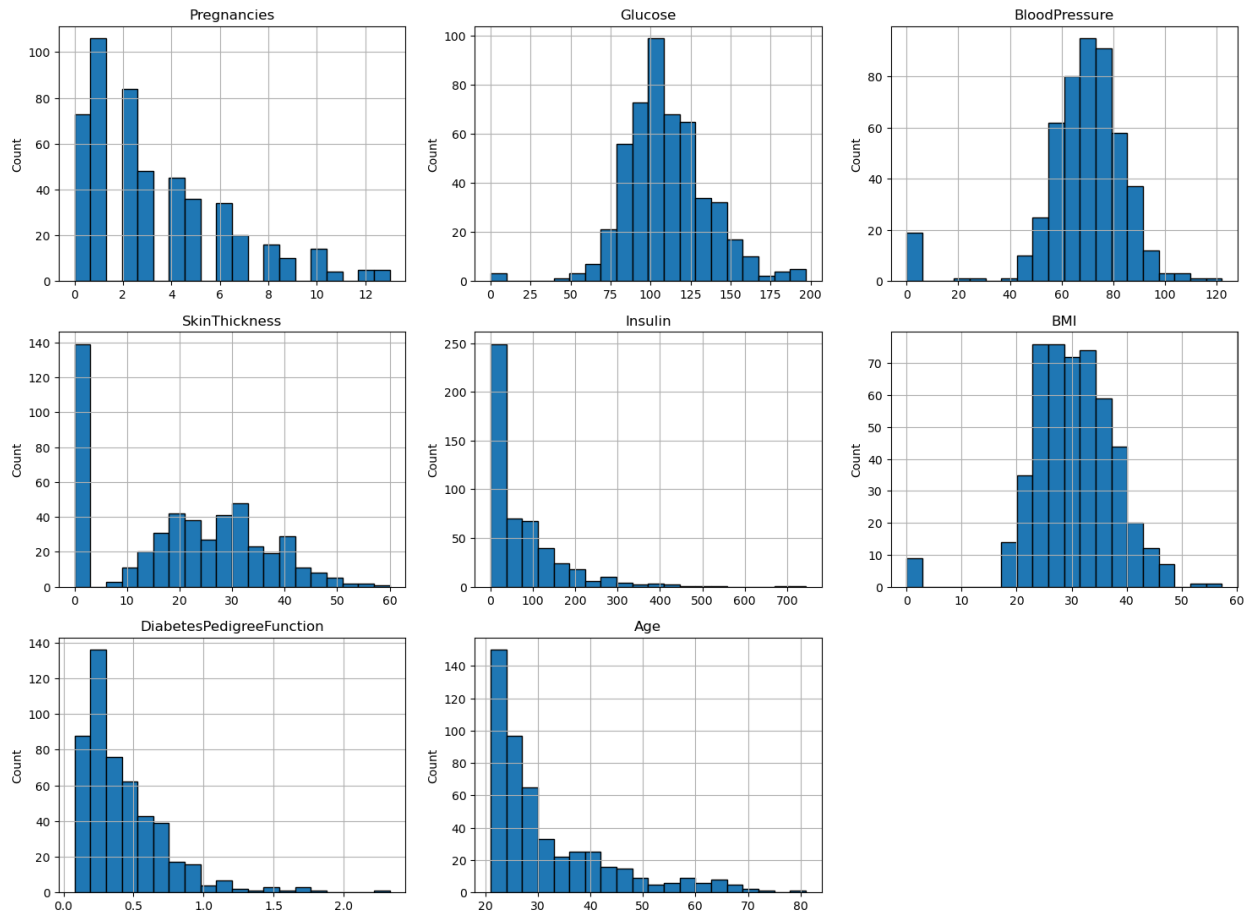
#General Analysis
data1 = df[df["Outcome"]==1]
columns = df.columns[:8]
plt.subplots(figsize=(18,18))
length=len(columns)
for i, j in zip(columns, range(length)):
    plt.subplot((length//2),3,j+1)
    plt.ylabel("Count")
    data1[i].hist(bins=20,edgecolor='black')
    plt.title(i)
plt.show()

```



Analysis of Non-Diabetic Cases

```
#General Analysis
data1 = df[df["Outcome"]==0]
columns = df.columns[:8]
plt.subplots(figsize=(18,18))
length=len(columns)
for i, j in zip(columns, range(length)):
    plt.subplot((length//2),3,j+1)
    plt.ylabel("Count")
    data1[i].hist(bins=20,edgecolor='black')
    plt.title(i)
plt.show()
```



```
#Split Data
```

```
p = df[df.Outcome == 1]
```

```
n = df[df.Outcome == 0]
```

```
plt.scatter(p.Pregnancies,p.Glucose,color = "brown",label="Diabet  
Positive",alpha=0.6)
```

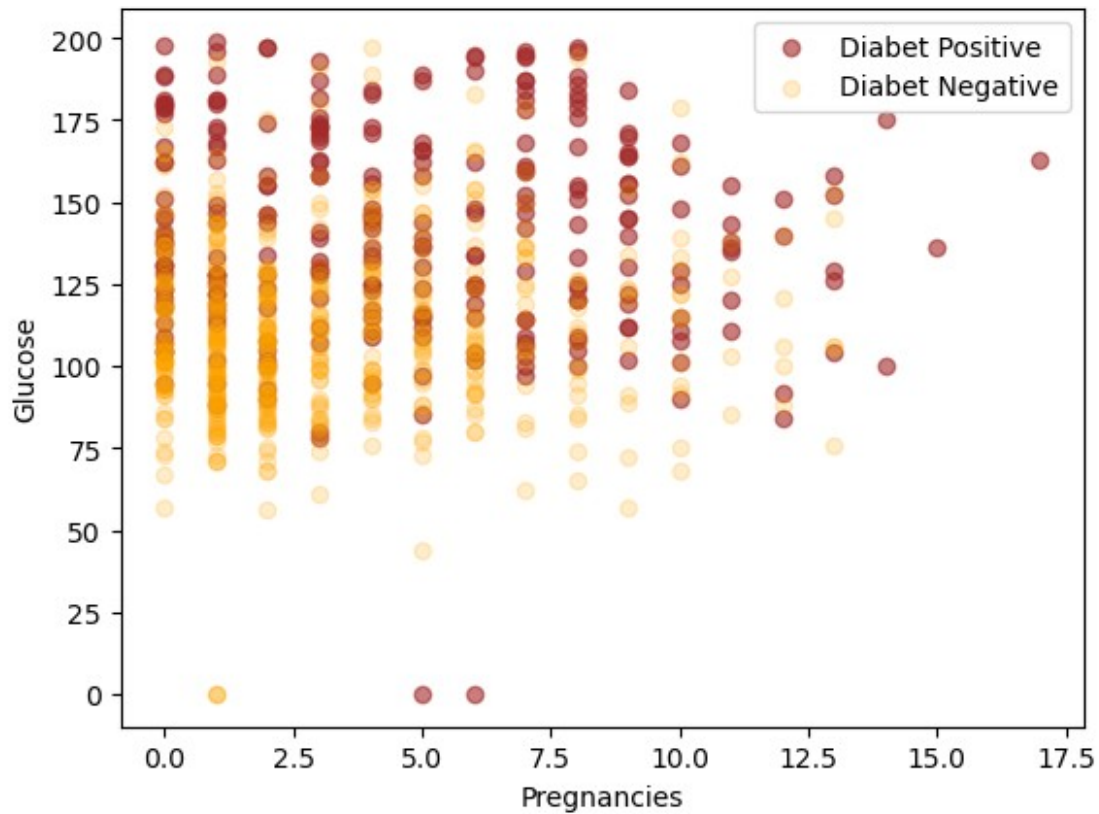
```
plt.scatter(n.Pregnancies,n.Glucose,color = "Orange",label="Diabet  
Negative",alpha=0.2)
```

```
plt.xlabel("Pregnancies")
```

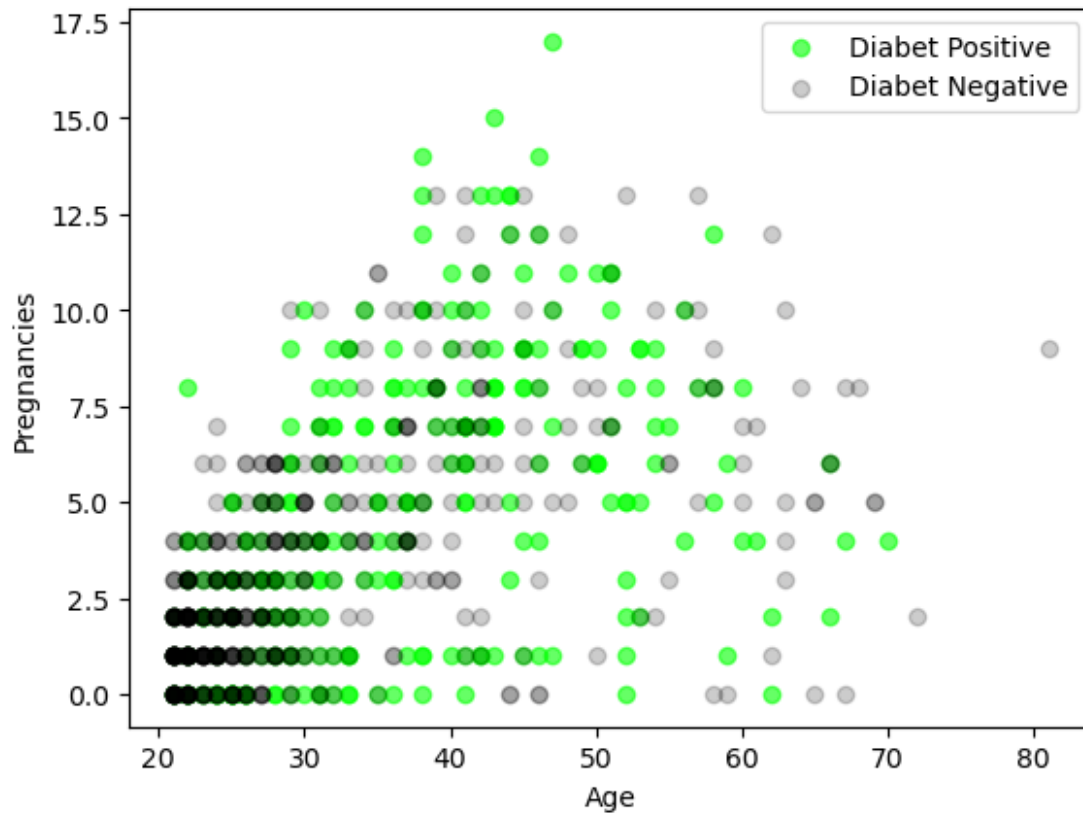
```
plt.ylabel("Glucose")
```

```
plt.legend()
```

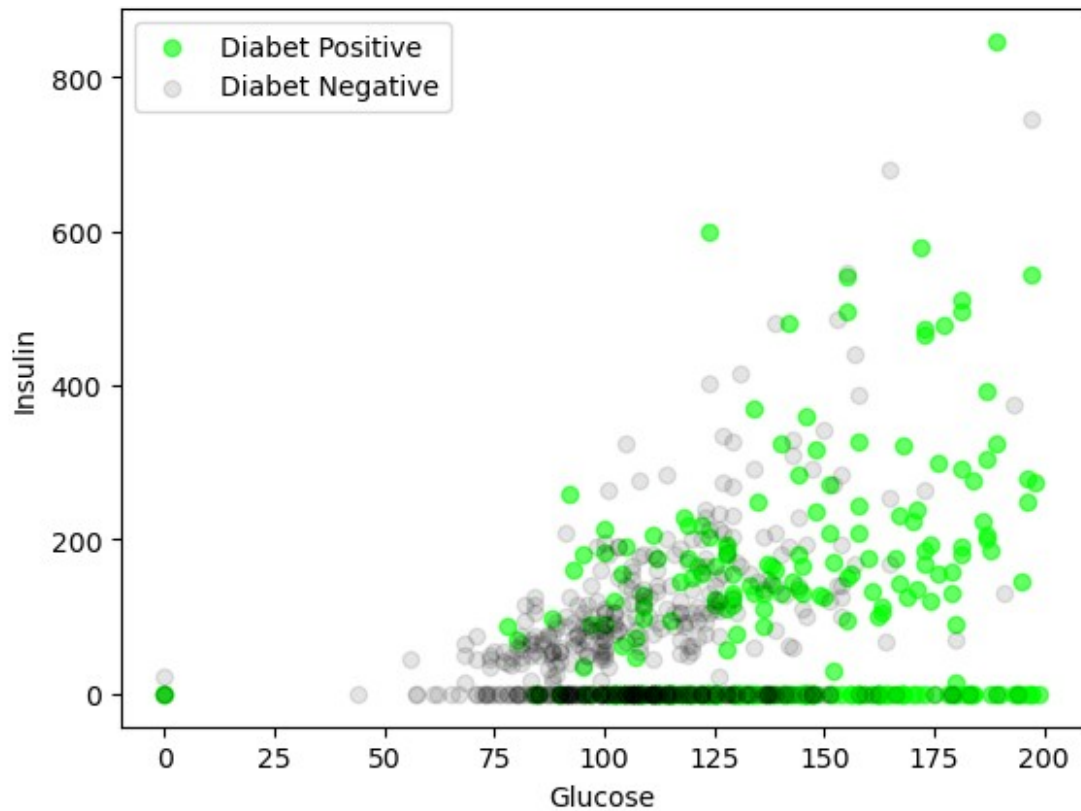
```
plt.show()
```



```
plt.scatter(p.Age,p.Pregnancies,color = "lime",label="Diabet  
Positive",alpha=0.6)  
plt.scatter(n.Age,n.Pregnancies,color = "black",label="Diabet  
Negative",alpha=0.2)  
plt.xlabel("Age")  
plt.ylabel("Pregnancies")  
plt.legend()  
plt.show()
```


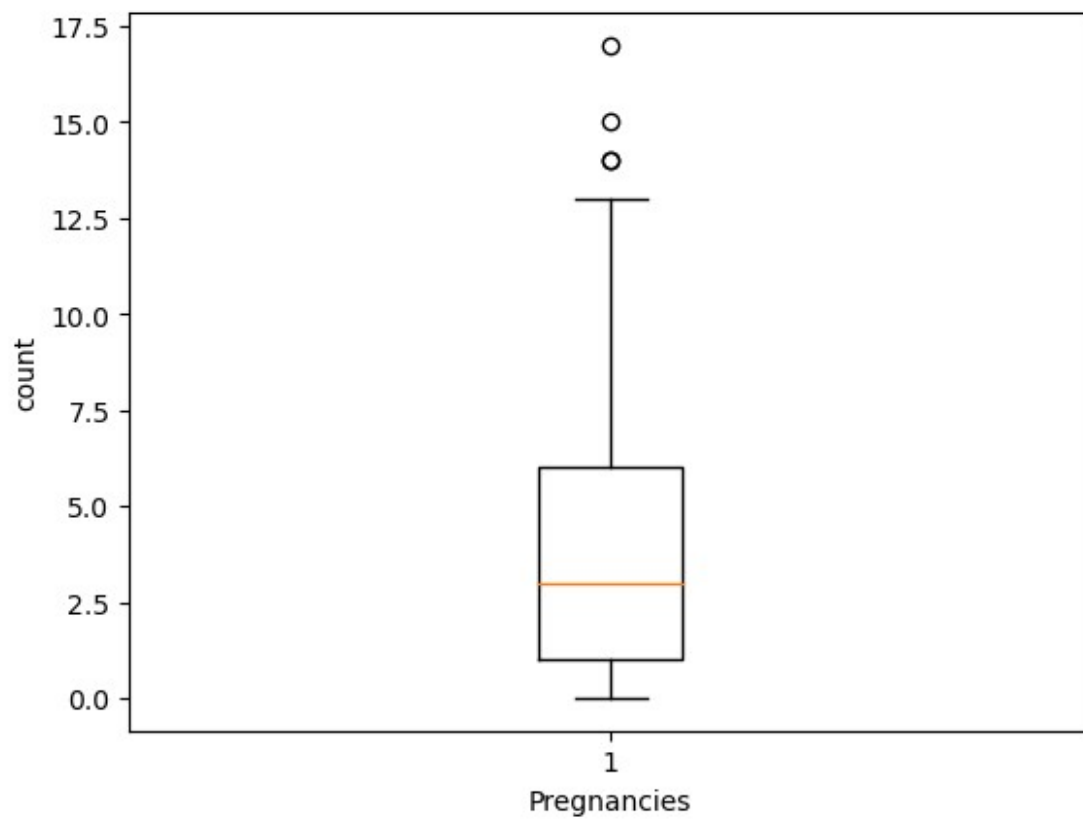


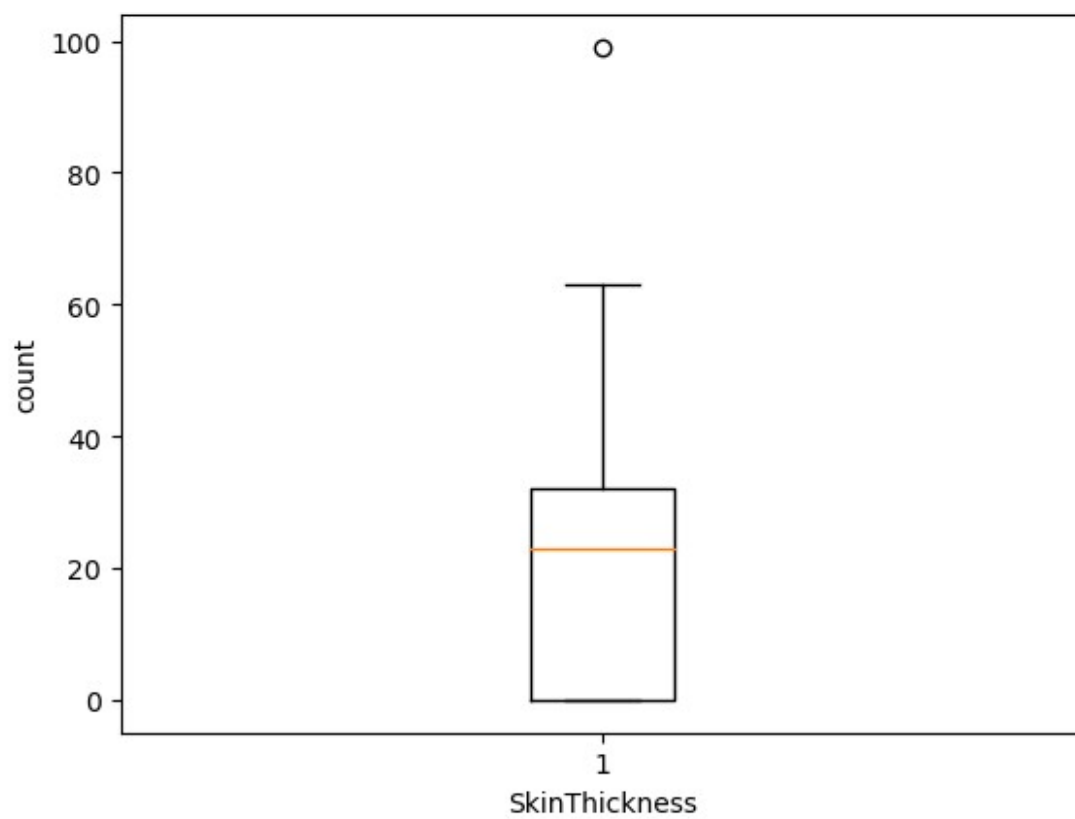
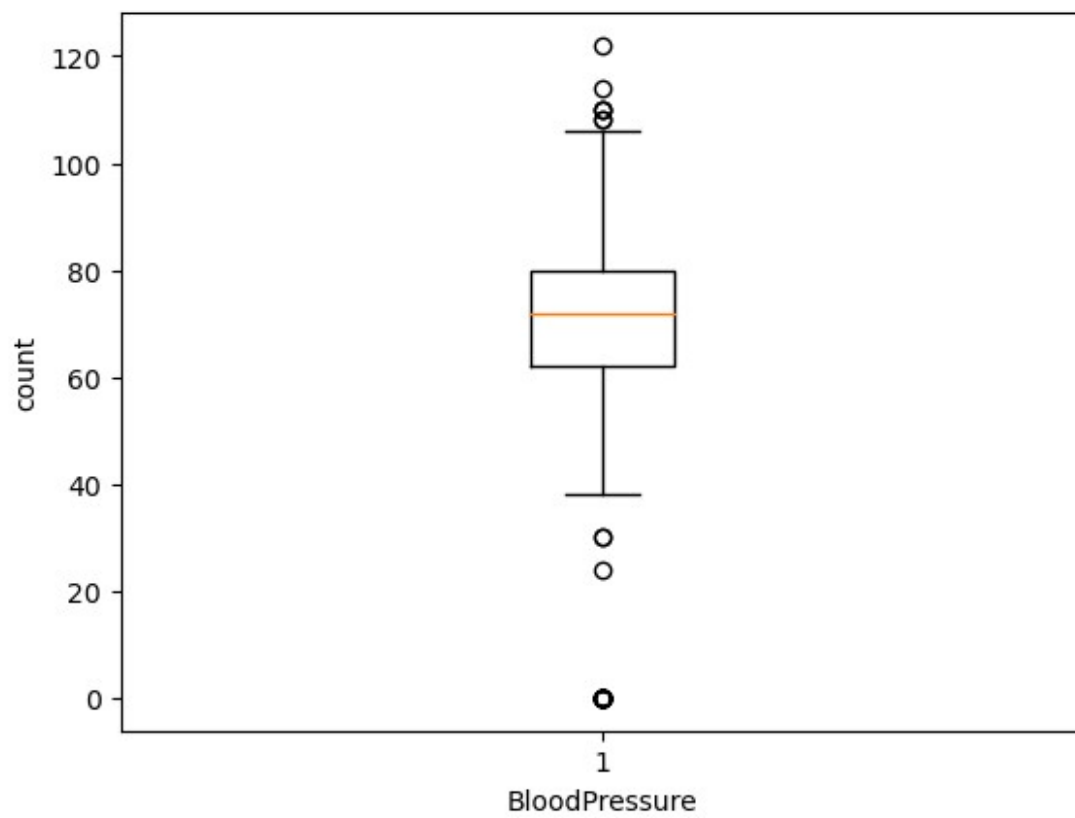
```
plt.scatter(p.Glucose,p.Insulin,color = "lime",label="Diabet
Positive",alpha=0.6)
plt.scatter(n.Glucose,n.Insulin,color = "black",label="Diabet
Negative",alpha=0.1)
plt.xlabel("Glucose")
plt.ylabel("Insulin")
plt.legend()
plt.show()
```

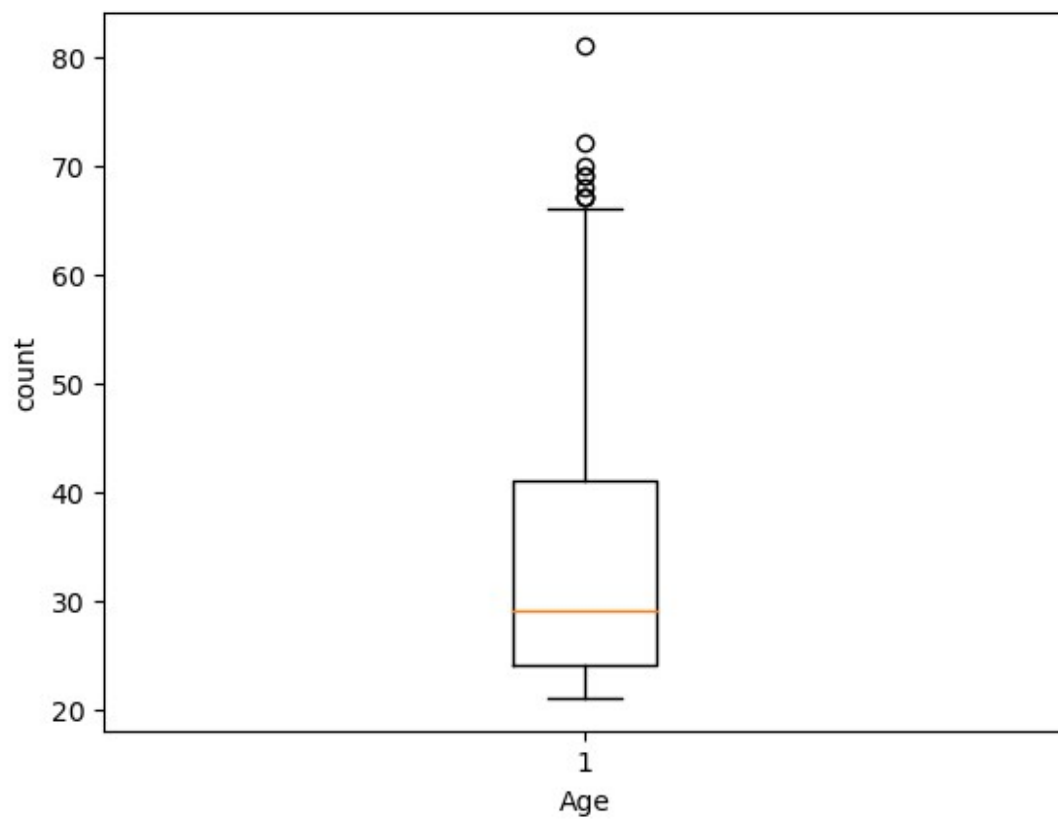
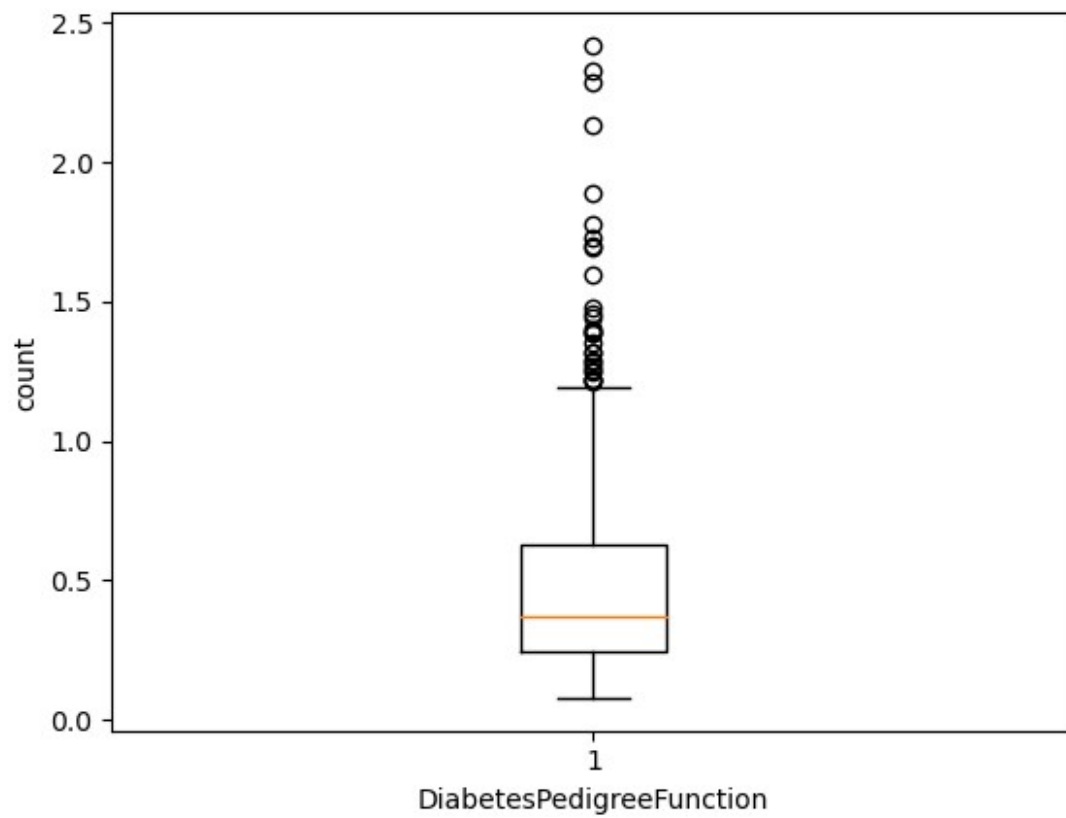


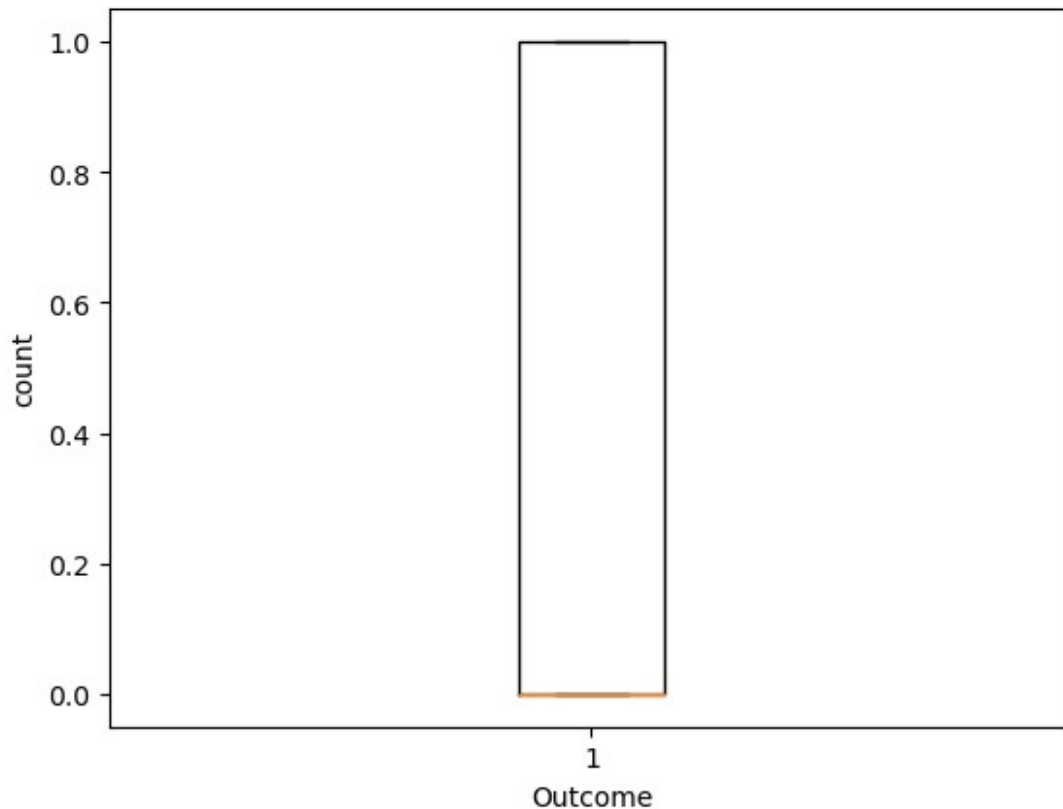
Identifying and Removal of Outlier

```
for col_name in df.columns:
    if(df[col_name].dtypes=='int64' or
df[col_name].dtypes=='float64'):
        plt.boxplot(df[col_name])
        plt.xlabel(col_name)
        plt.ylabel('count')
        plt.show()
```

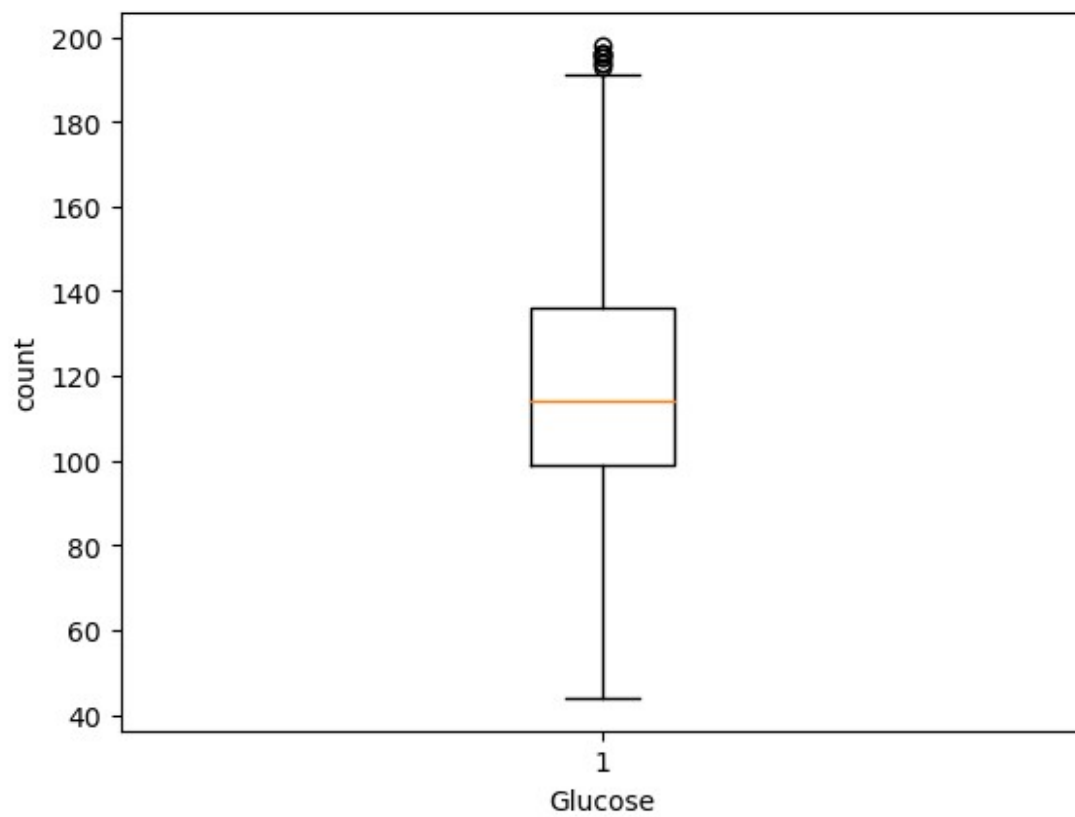
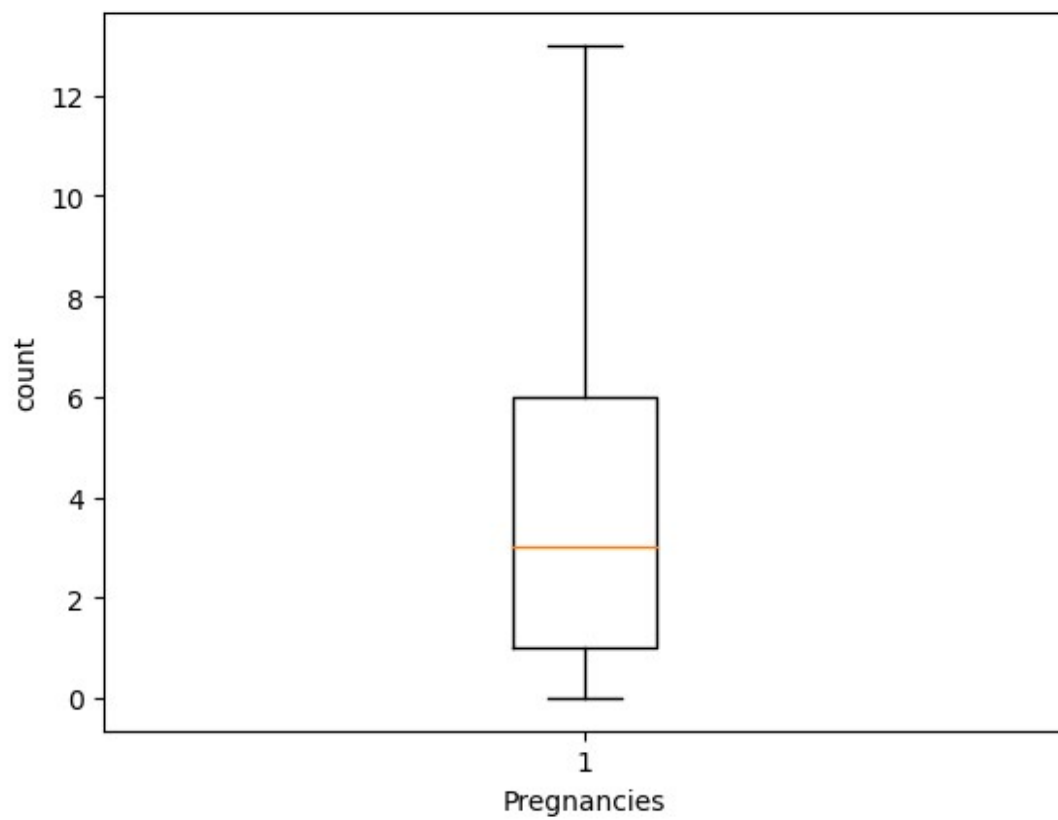


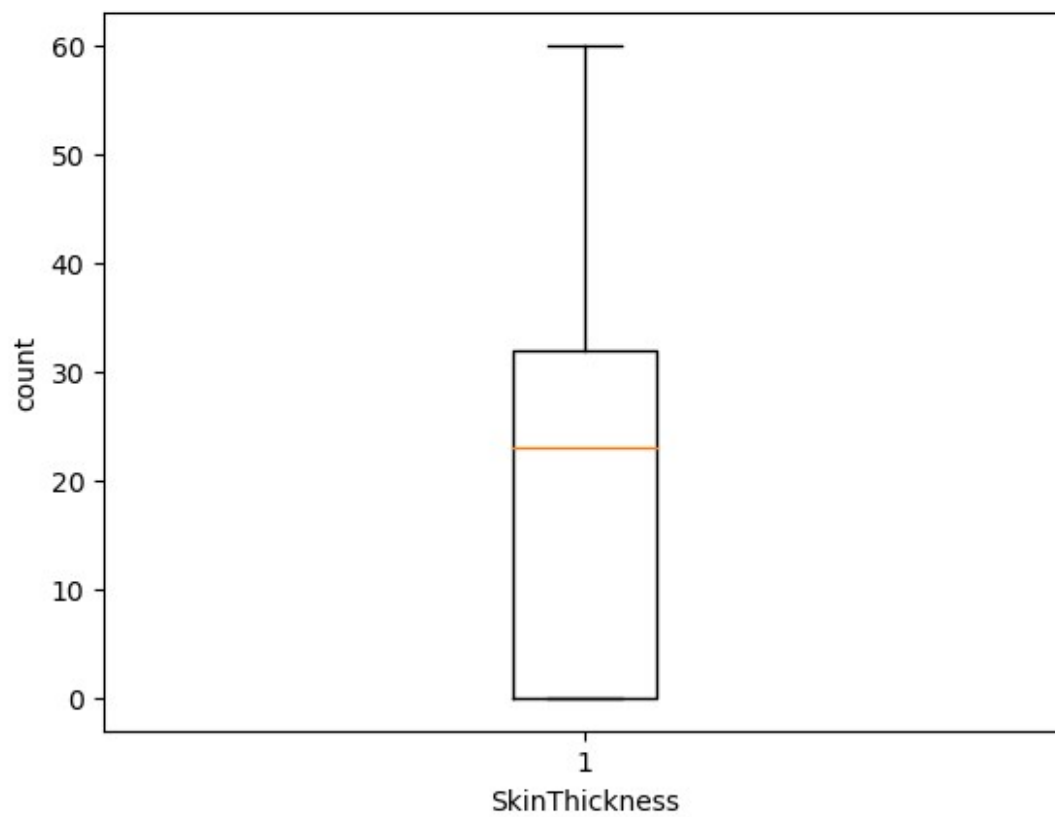
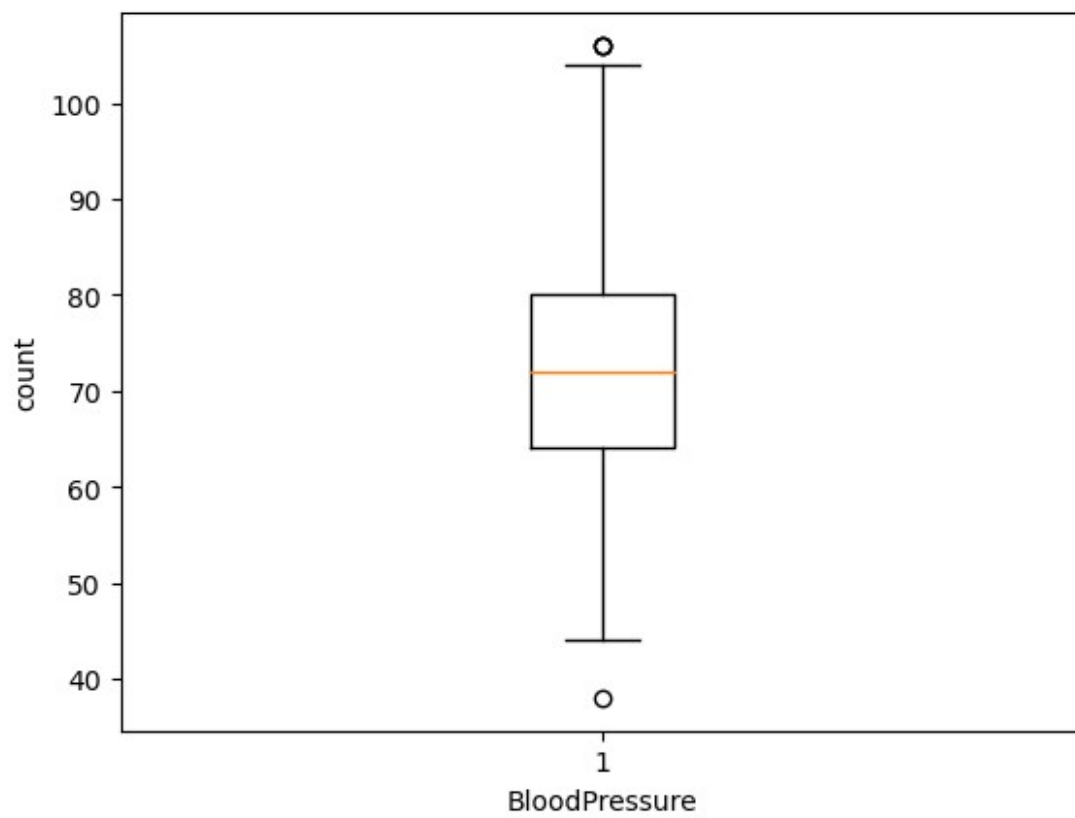



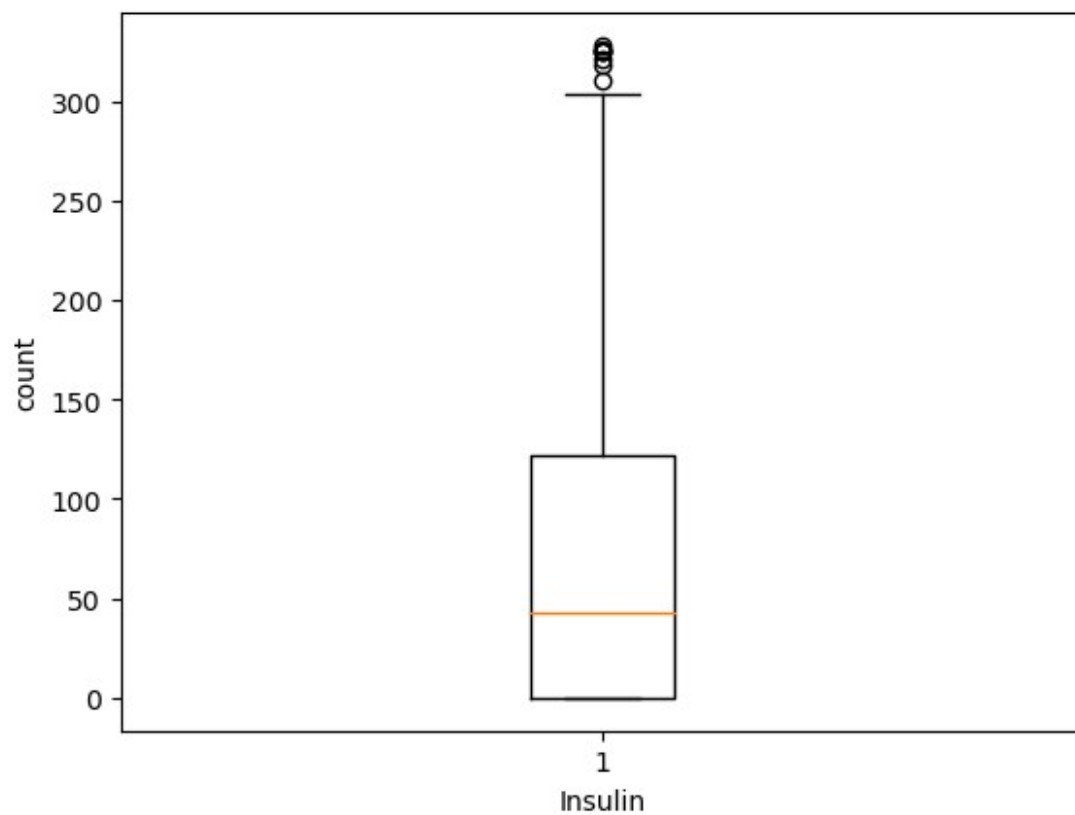


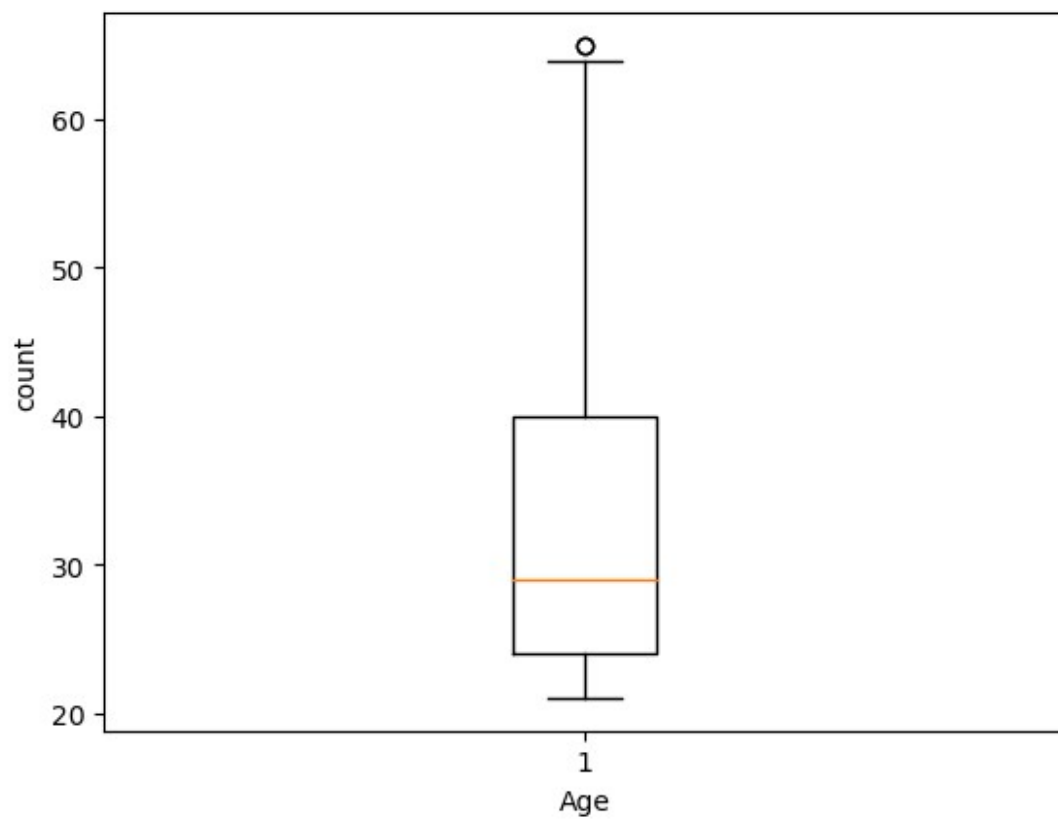
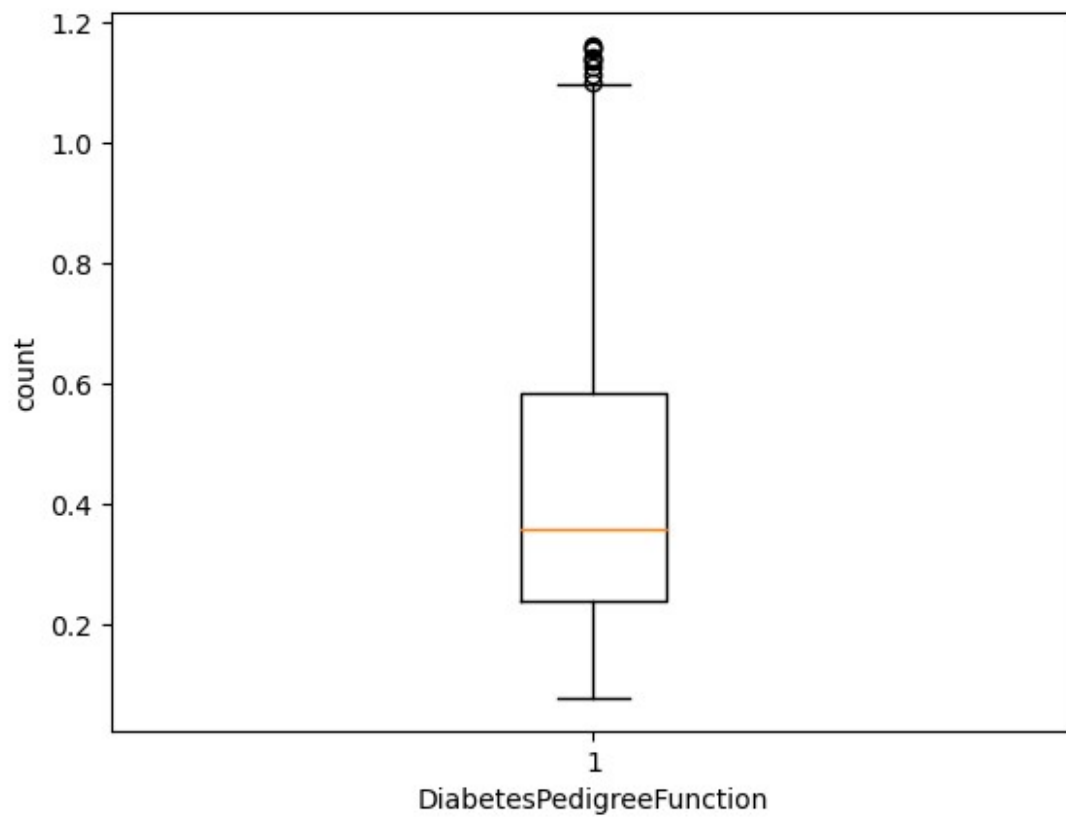
```
for i in df.columns:
    if(i!='Outcome'):
        q1=df[i].quantile(0.25)
        q3=df[i].quantile(0.75)
        iqr=q3-q1
        df=df[(df[i]>q1-1.5*iqr) & (df[i]<q3+1.5*iqr)]

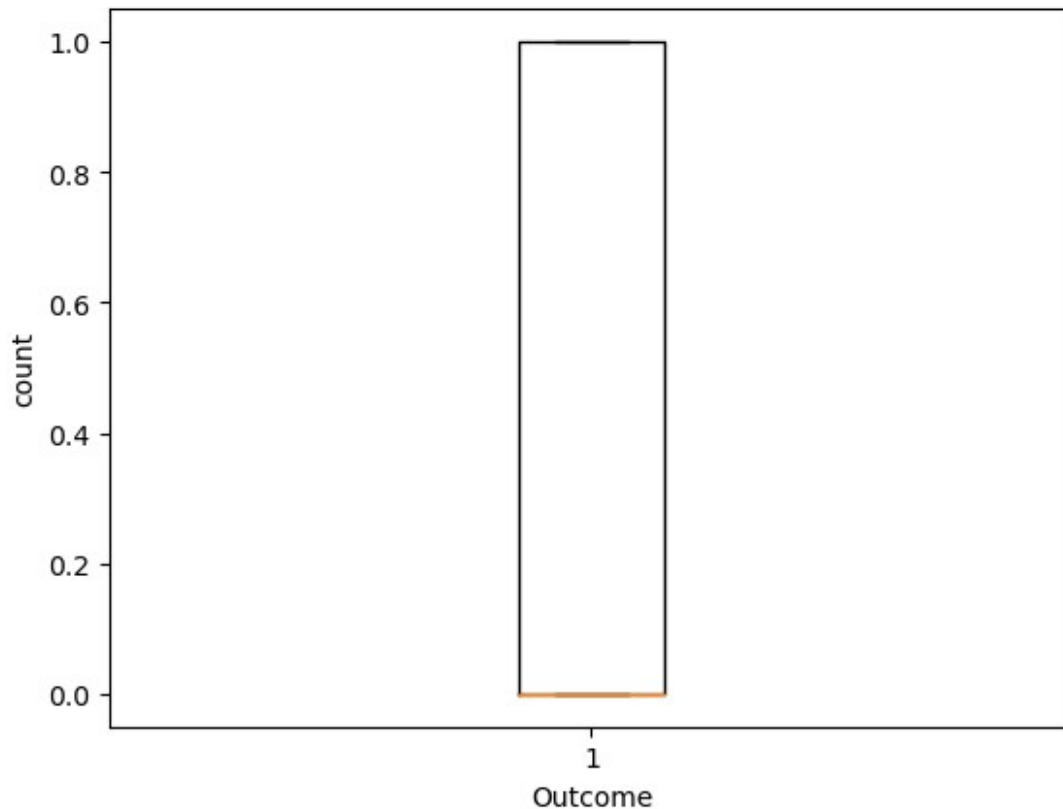
for col_name in df.columns:
    if(df[col_name].dtypes=='int64' or
df[col_name].dtypes=='float64'):
        plt.boxplot(df[col_name])
        plt.xlabel(col_name)
        plt.ylabel('count')
        plt.show()
```











Segregation of Data (Independent and Dependent)

```
x=df.drop('Outcome',axis=1)
y=df['Outcome']
```

Splitting of Dataset into Train and Test Data

```
test_size = 0.2

x_train, x_test, y_train, y_test =
train_test_split(x,y,test_size=test_size,random_state=42)
```

- StandardScaler is a preprocessing technique used in machine learning and statistics to scale (standardize) the features of a dataset. The goal of StandardScaler is to transform the data such that it has a mean of 0 and a standard deviation of 1. This process is also known as "z-score normalization" or "z-score scaling."

$Z = (X - \mu) / \sigma$ "Z" is the standardized value. "X" is the original value. " μ " is the mean of the feature. " σ " is the standard deviation of the feature.

- StandardScaler is a good choice when you have features with different scales or units, and you want to ensure that they all have the same scale for modeling purposes. It's commonly used with algorithms like Support Vector Machines (SVM), Principal Component Analysis (PCA), and k-Nearest Neighbors (k-NN).

```
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

Decision Tree Model

```
model = tree.DecisionTreeClassifier()
model.fit(x_train, y_train)

DecisionTreeClassifier()

prediction=model.predict(x_test)

print(metrics.classification_report(y_test,prediction))
```

	precision	recall	f1-score	support
0	0.86	0.76	0.81	97
1	0.45	0.61	0.52	31
accuracy			0.73	128
macro avg	0.66	0.69	0.66	128
weighted avg	0.76	0.73	0.74	128

```
print(metrics.accuracy_score(y_test,prediction))

0.7265625
```

Hyper Parameter Tuning

```
parameter={
    'max_depth':[None,10,20,30],
    'min_samples_leaf':[1,2,4],
    'min_samples_split':[2,5,10],
    'max_features':['auto','sqrt'],
    'criterion': ['gini', 'entropy']
}

rf_classifier = DecisionTreeClassifier(random_state=42)
grid_search=GridSearchCV( estimator=rf_classifier ,param_grid=parameter,
r,verbose=2,cv=5,n_jobs=-1)
grid_search.fit(x, y)
```

```
print("Best Parameters",grid_search.best_params_)
```

Fitting 5 folds for each of 144 candidates, totalling 720 fits
Best Parameters {'criterion': 'gini', 'max_depth': None,
'max_features': 'sqrt', 'min_samples_leaf': 4, 'min_samples_split': 2}

C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection_validation.py:425: FitFailedWarning:
360 fits failed out of a total of 720.
The score on these train-test partitions for these parameters will be
set to nan.
If these failures are not expected, you can try to debug them by
setting error_score='raise'.

Below are more details about the failures:

```
-----  
-----  
195 fits failed with the following error:  
Traceback (most recent call last):  
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\  
model_selection\_validation.py", line 729, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\  
base.py", line 1145, in wrapper  
    estimator._validate_params()  
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\  
base.py", line 638, in _validate_params  
    validate_parameter_constraints(  
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\  
utils\_param_validation.py", line 95, in  
validate_parameter_constraints  
    raise InvalidParameterError(  
sklearn.utils._param_validation.InvalidParameterError: The  
'max_features' parameter of DecisionTreeClassifier must be an int in  
the range [1, inf), a float in the range (0.0, 1.0], a str among  
{ 'sqrt', 'log2' } or None. Got 'auto' instead.
```

```
-----  
-----  
165 fits failed with the following error:  
Traceback (most recent call last):  
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\  
model_selection\_validation.py", line 729, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\  
base.py", line 1145, in wrapper  
    estimator._validate_params()  
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\  
base.py", line 638, in _validate_params
```

```

    validate_parameter_constraints(
        File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
utils\_param_validation.py", line 95, in
validate_parameter_constraints
        raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of DecisionTreeClassifier must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'log2', 'sqrt'} or None. Got 'auto' instead.

    warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_search.py:979: UserWarning: One or more of the test
scores are non-finite: [      nan      nan      nan      nan
nan      nan
      nan      nan      nan      nan      nan      nan
0.69811762 0.68235728 0.72956447 0.74685039 0.74685039 0.72962598
      nan      nan      nan      nan      nan      nan
      nan      nan      nan      nan      nan      nan
0.73897638 0.71226624 0.72641486 0.74533711 0.74533711 0.73277559
      nan      nan      nan      nan      nan      nan
      nan      nan      nan      nan      nan      nan
0.69811762 0.68235728 0.72956447 0.74685039 0.74685039 0.72962598
      nan      nan      nan      nan      nan      nan
      nan      nan      nan      nan      nan      nan
0.69811762 0.68235728 0.72956447 0.74685039 0.74685039 0.72962598
      nan      nan      nan      nan      nan      nan
      nan      nan      nan      nan      nan      nan
0.72331447 0.71703986 0.70441683 0.72647638 0.72647638 0.73582677
      nan      nan      nan      nan      nan      nan
      nan      nan      nan      nan      nan      nan
0.71385335 0.71384104 0.71686762 0.72645177 0.72645177 0.73115157
      nan      nan      nan      nan      nan      nan
      nan      nan      nan      nan      nan      nan
0.72331447 0.71703986 0.70441683 0.72647638 0.72647638 0.73582677
      nan      nan      nan      nan      nan      nan
      nan      nan      nan      nan      nan      nan
0.72331447 0.71703986 0.70441683 0.72647638 0.72647638 0.73582677]
    warnings.warn(

model = tree.DecisionTreeClassifier(criterion= 'gini', max_depth=
None, max_features= 'sqrt', min_samples_leaf= 4, min_samples_split= 2)
model.fit(x_train, y_train)

DecisionTreeClassifier(max_features='sqrt', min_samples_leaf=4)

prediction=model.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test,prediction))

Accuracy: 0.671875

```

Random Forest Model

```
model=RandomForestClassifier()  
model.fit(x_train,y_train)  
prediction=model.predict(x_test)  
  
print(metrics.classification_report(y_test,prediction))
```

	precision	recall	f1-score	support
0	0.86	0.85	0.85	97
1	0.55	0.58	0.56	31
accuracy			0.78	128
macro avg	0.70	0.71	0.71	128
weighted avg	0.79	0.78	0.78	128

```
print(metrics.accuracy_score(y_test,prediction))  
  
0.78125
```

Hyper Parameter Tuning

```
parameter={  
    'max_depth':[None,10,20,30],  
    'min_samples_leaf':[1,2,4],  
    'min_samples_split':[2,5,10],  
    'max_features':['auto','sqrt'],  
    'n_estimators':[100,200,300]  
}  
rf_classifier = RandomForestClassifier(random_state=42)  
grid_search=GridSearchCV( estimator=rf_classifier ,param_grid=paramete  
r,verbose=2,cv=5,n_jobs=-1)  
grid_search.fit(x, y)
```

Fitting 5 folds for each of 216 candidates, totalling 1080 fits

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\  
model_selection\_validation.py:425: FitFailedWarning:  
540 fits failed out of a total of 1080.  
The score on these train-test partitions for these parameters will be  
set to nan.  
If these failures are not expected, you can try to debug them by  
setting error_score='raise'.
```

Below are more details about the failures:

```
-----  
-----  
158 fits failed with the following error:
```

```

Traceback (most recent call last):
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_validation.py", line 729, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 1145, in wrapper
    estimator._validate_params()
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 638, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of RandomForestClassifier must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'sqrt', 'log2'} or None. Got 'auto' instead.

```

382 fits failed with the following error:

```

Traceback (most recent call last):
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_validation.py", line 729, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 1145, in wrapper
    estimator._validate_params()
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
base.py", line 638, in _validate_params
    validate_parameter_constraints(
  File "C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of RandomForestClassifier must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'log2', 'sqrt'} or None. Got 'auto' instead.

```

```

    warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\
model_selection\_search.py:979: UserWarning: One or more of the test
scores are non-finite: [      nan      nan      nan      nan
nan      nan
      nan      nan      nan      nan      nan      nan      nan
      nan      nan      nan      nan      nan      nan      nan
      nan      nan      nan      nan      nan      nan      nan

```

```

nan nan nan 0.78458415 0.77672244 0.79243356
0.78302165 0.78617126 0.78302165 0.78930856 0.78617126 0.78458415
0.78613435 0.78457185 0.78617126 0.78142224 0.78300935 0.78614665
0.78460876 0.78930856 0.78930856 0.77989665 0.78615896 0.78772146
0.77989665 0.78615896 0.78772146 0.78300935 0.78773376 0.79243356
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan 0.79558317 0.78615896 0.78930856
0.78459646 0.79089567 0.78932087 0.78932087 0.78615896 0.78615896
0.78300935 0.78145915 0.78617126 0.78929626 0.78458415 0.78773376
0.78144685 0.79087106 0.78458415 0.78300935 0.78615896 0.78930856
0.78300935 0.78615896 0.78930856 0.78302165 0.78615896 0.79088337
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan 0.78458415 0.77672244 0.79243356
0.78302165 0.78617126 0.78302165 0.78930856 0.78617126 0.78458415
0.78613435 0.78457185 0.78617126 0.78142224 0.78300935 0.78614665
0.78460876 0.78930856 0.78930856 0.77989665 0.78615896 0.78772146
0.77989665 0.78615896 0.78772146 0.78300935 0.78773376 0.79243356
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan nan nan nan
nan nan nan 0.78458415 0.77672244 0.79243356
0.78302165 0.78617126 0.78302165 0.78930856 0.78617126 0.78458415
0.78613435 0.78457185 0.78617126 0.78142224 0.78300935 0.78614665
0.78460876 0.78930856 0.78930856 0.77989665 0.78615896 0.78772146
0.77989665 0.78615896 0.78772146 0.78300935 0.78773376 0.79243356]
warnings.warn(

```

```

GridSearchCV(cv=5, estimator=RandomForestClassifier(random_state=42),
n_jobs=-1,

```

```

    param_grid={'max_depth': [None, 10, 20, 30],
                'max_features': ['auto', 'sqrt'],
                'min_samples_leaf': [1, 2, 4],
                'min_samples_split': [2, 5, 10],
                'n_estimators': [100, 200, 300]},
    verbose=2)

```

```

print("Best Hyperparameters:", grid_search.best_params_)

```

```

Best Hyperparameters: {'max_depth': 10, 'max_features': 'sqrt',
'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}

```

```

model=RandomForestClassifier(max_depth=10,max_features='sqrt',min_samp
les_leaf=1,min_samples_split=2,n_estimators=100)

```



```

model.fit(x_train,y_train)
prediction=model.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test,prediction))

```

Accuracy: 0.765625

KNN Model

```

model=KNeighborsClassifier(n_neighbors=3)
model.fit(x_train,y_train)
prediction=model.predict(x_test)

print(metrics.classification_report(y_test,prediction))

```

	precision	recall	f1-score	support
0	0.86	0.84	0.85	97
1	0.53	0.58	0.55	31
accuracy			0.77	128
macro avg	0.70	0.71	0.70	128
weighted avg	0.78	0.77	0.78	128

```

print(metrics.accuracy_score(y_test,prediction))

```

0.7734375

Hyper Parameter Tuning

```

# Define the hyperparameter grid
param_grid = {
    'n_neighbors': [3,5,7,9,11], # Example values for n_neighbors
    'weights': ['uniform', 'distance'], # Weighting options
    'metric': ['euclidean', 'manhattan'] # Distance metrics
}

# Create the KNN classifier
knn = KNeighborsClassifier()

# Perform grid search with cross-validation
grid_search = GridSearchCV(estimator=knn, param_grid=param_grid,
    scoring='accuracy', cv=5)
grid_search.fit(x_train, y_train)
print("best parameters :",grid_search.best_params_)

best parameters : {'metric': 'euclidean', 'n_neighbors': 11,
'weights': 'distance'}

```

```
hyper_knn =
KNeighborsClassifier(n_neighbors=11,metric="euclidean",weights="distance")
hyper_knn.fit(x_train,y_train)
hyper_prediction = hyper_knn.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test,hyper_prediction))
```

Accuracy: 0.734375

Naive Bayes Model

```
from sklearn.naive_bayes import GaussianNB
model=GaussianNB()
model.fit(x_train,y_train)
prediction=model.predict(x_test)

print(metrics.classification_report(y_test,prediction))
```

	precision	recall	f1-score	support
0	0.83	0.79	0.81	97
1	0.43	0.48	0.45	31
accuracy			0.72	128
macro avg	0.63	0.64	0.63	128
weighted avg	0.73	0.72	0.72	128

```
print(metrics.accuracy_score(y_test,prediction))
```

0.71875

Logistic Regression Model

```
logistic=LogisticRegression(solver="liblinear")
logistic_model=logistic.fit(x,y)
logistic_model

LogisticRegression(solver='liblinear')

prediction=logistic.predict(x_test)
```

C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\base.py:465: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
warnings.warn(

```
print(metrics.classification_report(y_test,prediction))
```

	precision	recall	f1-score	support
0	0.76	1.00	0.86	97
1	0.00	0.00	0.00	31
accuracy			0.76	128
macro avg	0.38	0.50	0.43	128
weighted avg	0.57	0.76	0.65	128

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\metrics\
_classification.py:1469: UndefinedMetricWarning: Precision and F-score
are ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\metrics\
_classification.py:1469: UndefinedMetricWarning: Precision and F-score
are ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\sklearn\metrics\
_classification.py:1469: UndefinedMetricWarning: Precision and F-score
are ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
print("Accuracy:",metrics.accuracy_score(y_test,prediction))
```

```
Accuracy: 0.7578125
```

Conclusion

- Random forest Classifier gives more accuracy than other model