

Importing Packages

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Data Preprocessing

```
df=pd.read_csv(r"C:\Users\Arigala.Adarsh\Downloads\train (23).csv")
```

```
df.head()
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
0	1000001	P00069042	F	0-17	10.0	A	
1	1000001	P00248942	F	0-17	10.0	A	
2	1000001	P00087842	F	0-17	10.0	A	
3	1000001	P00085442	F	0-17	10.0	A	
4	1000002	P00285442	M	55+	16.0	C	

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
0	2	0.0	3.0	
1	2	0.0	1.0	
2	2	0.0	12.0	
3	2	0.0	12.0	
4	4+	0.0	8.0	

	Product_Category_2	Product_Category_3	Purchase
0	NaN	NaN	8370.0
1	6.0	14.0	15200.0
2	NaN	NaN	1422.0
3	14.0	NaN	1057.0
4	NaN	NaN	7969.0

```
df.shape
```

```
(263015, 12)
```

Exploratory Data Analysis(EDA)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 263015 entries, 0 to 263014
Data columns (total 12 columns):
```

#	Column	Non-Null Count		Dtype
---	-----	-----		-----
0	User_ID	263015	non-null	int64
1	Product_ID	263014	non-null	object
2	Gender	263014	non-null	object
3	Age	263014	non-null	object
4	Occupation	263014	non-null	float64
5	City_Category	263014	non-null	object
6	Stay_In_Current_City_Years	263014	non-null	object
7	Marital_Status	263014	non-null	float64
8	Product_Category_1	263014	non-null	float64
9	Product_Category_2	181501	non-null	float64
10	Product_Category_3	80582	non-null	float64
11	Purchase	263014	non-null	float64

dtypes: float64(6), int64(1), object(5)
memory usage: 24.1+ MB

df.dtypes

User_ID	int64
Product_ID	object
Gender	object
Age	object
Occupation	float64
City_Category	object
Stay_In_Current_City_Years	object
Marital_Status	float64
Product_Category_1	float64
Product_Category_2	float64
Product_Category_3	float64
Purchase	float64

dtype: object

df.describe()

	User_ID	Occupation	Marital_Status	Product_Category_1
\				
count	2.630150e+05	263014.000000	263014.000000	263014.000000
mean	1.002941e+06	8.083558	0.408685	5.291099
std	2.593126e+03	6.524052	0.491592	3.745722
min	1.000000e+01	0.000000	0.000000	1.000000
25%	1.001457e+06	2.000000	0.000000	1.000000
50%	1.002972e+06	7.000000	0.000000	5.000000
75%	1.004335e+06	14.000000	1.000000	8.000000

max	1.006040e+06	20.000000	1.000000	18.000000
-----	--------------	-----------	----------	-----------

	Product_Category_2	Product_Category_3	Purchase
count	181501.000000	80582.000000	263014.000000
mean	9.844756	12.658298	9319.305269
std	5.086696	4.129156	4970.152966
min	2.000000	3.000000	185.000000
25%	5.000000	9.000000	5863.000000
50%	9.000000	14.000000	8060.000000
75%	15.000000	16.000000	12059.000000
max	18.000000	18.000000	23961.000000

```
df.isnull().sum()
```

User_ID	0
Product_ID	1
Gender	1
Age	1
Occupation	1
City_Category	1
Stay_In_Current_City_Years	1
Marital_Status	1
Product_Category_1	1
Product_Category_2	81514
Product_Category_3	182433
Purchase	1
dtype: int64	

```
df.dropna(inplace=True)
```

```
df.shape
```

```
(80582, 12)
```

```
df.columns
```

```
Index(['User_ID', 'Product_ID', 'Gender', 'Age', 'Occupation',  
      'City_Category',  
      'Stay_In_Current_City_Years', 'Marital_Status',  
      'Product_Category_1',  
      'Product_Category_2', 'Product_Category_3', 'Purchase'],  
      dtype='object')
```

```
from sklearn.preprocessing import LabelEncoder  
le=LabelEncoder()  
df['User_ID'] = le.fit_transform(df['User_ID'])  
df['Product_ID'] = le.fit_transform(df['Product_ID'])  
df['Gender'] = le.fit_transform(df['Gender'])  
df['Age'] = le.fit_transform(df['Age'])
```

```
df['City_Category'] = le.fit_transform(df['City_Category'])
```

```
df.dtypes
```

```
User_ID          int64
Product_ID       int32
Gender           int32
Age             int32
Occupation       float64
City_Category    int32
Stay_In_Current_City_Years  object
Marital_Status   float64
Product_Category_1 float64
Product_Category_2 float64
Product_Category_3 float64
Purchase         float64
dtype: object
```

```
df
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
1	0	391	0	0	10.0	0	
6	3	284	1	4	7.0	1	
13	4	211	1	2	20.0	0	
14	5	363	0	5	9.0	0	
16	5	517	0	5	9.0	0	
...	
262997	4232	402	0	4	16.0	1	
263001	4232	337	0	4	16.0	1	
263003	4232	480	0	4	16.0	1	
263006	4232	49	0	4	16.0	1	
263011	4233	159	1	3	1.0	1	

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
1	2	0.0	1.0	
6	2	1.0	1.0	
13	1	1.0	1.0	
14	1	0.0	5.0	
16	1	0.0	2.0	
...	
262997	0	1.0	1.0	
263001	0	1.0	1.0	

263003	0	1.0	1.0
263006	0	1.0	8.0
263011	3	0.0	1.0

	Product_Category_2	Product_Category_3	Purchase
1	6.0	14.0	15200.0
6	8.0	17.0	19215.0
13	2.0	5.0	15665.0
14	8.0	14.0	5378.0
16	3.0	4.0	13055.0
...
262997	11.0	16.0	15175.0
263001	11.0	15.0	15430.0
263003	2.0	15.0	15387.0
263006	13.0	16.0	5861.0
263011	8.0	17.0	19253.0

[80582 rows x 12 columns]

df[df["Purchase"]==19060]

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
301	48	2	1	1	4.0	2	
18116	2695	71	1	2	0.0	1	
30806	4448	145	1	1	20.0	2	
34234	4980	149	1	1	4.0	2	
63488	3547	149	0	1	1.0	1	
99973	3305	107	1	2	3.0	0	
106330	4139	371	0	2	5.0	2	
113162	5157	145	1	3	1.0	2	
118000	192	138	1	1	4.0	1	
122437	874	284	0	6	16.0	2	
130088	1871	33	1	2	14.0	0	
163877	1237	159	1	4	7.0	1	
202902	1226	179	1	2	14.0	2	
216176	3214	2	1	2	17.0	2	

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1
301	0	0.0	1.0
18116	2	0.0	1.0
30806	1	0.0	1.0
34234	3	0.0	1.0

63488	4+	0.0	1.0
99973	2	0.0	1.0
106330	3	0.0	1.0
113162	0	1.0	1.0
118000	1	0.0	1.0
122437	4+	1.0	1.0
130088	3	0.0	1.0
163877	3	1.0	1.0
202902	1	1.0	1.0
216176	1	0.0	1.0

	Product_Category_2	Product_Category_3	Purchase
301	6.0	16.0	19060.0
18116	2.0	15.0	19060.0
30806	2.0	8.0	19060.0
34234	2.0	14.0	19060.0
63488	2.0	14.0	19060.0
99973	15.0	17.0	19060.0
106330	15.0	16.0	19060.0
113162	2.0	8.0	19060.0
118000	2.0	15.0	19060.0
122437	8.0	17.0	19060.0
130088	6.0	15.0	19060.0
163877	8.0	17.0	19060.0
202902	2.0	15.0	19060.0
216176	6.0	16.0	19060.0


```
df.isnull().sum()
```

User_ID	0
Product_ID	0
Gender	0
Age	0
Occupation	0
City_Category	0
Stay_In_Current_City_Years	0
Marital_Status	0
Product_Category_1	0
Product_Category_2	0
Product_Category_3	0

```

Purchase                                0
dtype: int64


df['Stay_In_Current_City_Years'].unique()
array(['2', '1', '4+', '0', '3'], dtype=object)

df['Stay_In_Current_City_Years']=df['Stay_In_Current_City_Years'].replace('4+', '4')

#changing the datatype from string to integer
df['Stay_In_Current_City_Years'] =
df['Stay_In_Current_City_Years'].astype(int)

sns.catplot(data=df, x='Gender', y='User_ID', kind='bar')
plt.show()

```




```

df.Gender.value_counts()
1    62549
0    18033
Name: Gender, dtype: int64

```


```
sns.catplot(data=df, x='Age', y='User_ID', hue='Gender', kind='bar',
height=6)
plt.show()
```



Inferences from the Plot:

1. The most frequent purchases came from the females having user ids that belong to age category 4(36-45)
2. The least frequent purchases came from the males of the age category 1(18-25), and females of age category 0(0-17).


```
sns.catplot(data=df, x='Marital_Status', y='User_ID', hue='Gender',
kind='bar', height=6)
plt.show()
```

the gender ratio in unmarried customers is almost similar, whereas the married customers have a slightly higher number of females in the lot.

```
plt.figure(figsize=(10,5))
sns.countplot(df['Product_Category_1'])
plt.show()
```


```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
_decorators.py:36: FutureWarning: Pass the following variable as a
keyword arg: x. From version 0.12, the only valid positional argument
will be `data`, and passing other arguments without an explicit
keyword will result in an error or misinterpretation.
  warnings.warn(
```



The product_category_1 sees a great rise of product category 1 and diminishes with the other products. The other considerable categories are 5, 2, 3, 6, 8, etc.

```
plt.figure(figsize=(10,5))
sns.countplot(df['Product_Category_2'])
plt.show()
```


```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
_decorators.py:36: FutureWarning: Pass the following variable as a
keyword arg: x. From version 0.12, the only valid positional argument
will be `data`, and passing other arguments without an explicit
keyword will result in an error or misinterpretation.
  warnings.warn(
```



Product_Category_2 sees a considerable balance among categories. With category 2 topping the charts, and other considerable categories are 8, 4, 5, 6, 14, 15, etc.

```
plt.figure(figsize=(10,5))
sns.countplot(df['Product_Category_3'])
plt.show()
```


```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
_decorators.py:36: FutureWarning: Pass the following variable as a
keyword arg: x. From version 0.12, the only valid positional argument
will be `data`, and passing other arguments without an explicit
keyword will result in an error or misinterpretation.
  warnings.warn(
```



For product_category_3, the balance is towards the end with 16 topping the charts, and other considerable categories are 15, 14, 5, 8, 9, 17, etc.

```
plt.figure(figsize=(10,5))
sns.countplot(df['Occupation'])
plt.show()
```


```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
_decorators.py:36: FutureWarning: Pass the following variable as a
keyword arg: x. From version 0.12, the only valid positional argument
will be `data`, and passing other arguments without an explicit
keyword will result in an error or misinterpretation.
  warnings.warn(
```




occupation also sees a constant balance with 4 topping the chart, and other categories in the considerable amount with 0,1,2,7,12,17,20.

```
plt.figure(figsize=(10,5))
sns.countplot(df['Stay_In_Current_City_Years'])
plt.show()
```

```
C:\Users\Arigala.Adarsh\anaconda3\lib\site-packages\seaborn\
_decorators.py:36: FutureWarning: Pass the following variable as a
keyword arg: x. From version 0.12, the only valid positional argument
will be `data`, and passing other arguments without an explicit
keyword will result in an error or misinterpretation.
  warnings.warn(
```



```
plt.figure(figsize=(18,7))
sns.heatmap(df.corr(),annot=True)
plt.show()
```



Statistical Analysis

A) It was observed that the average purchase made by the Men of the age 18-25 was 10000. Is it still the same?

- One Sample Test For Mean

```
new_data = df.loc[(df['Age'] == 1) & df['Gender'] == 1]
```

```
new_data.shape
```

```
(11904, 12)
```

```
new_data
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
98	20	487	1	1	15.0	0	
103	20	402	1	1	15.0	0	
111	20	358	1	1	15.0	0	
127	23	332	1	1	4.0	2	
128	23	71	1	1	4.0	2	
...	
262758	4206	180	1	1	18.0	1	
262759	4206	394	1	1	18.0	1	
262760	4206	23	1	1	18.0	1	
262764	4206	75	1	1	18.0	1	
262952	4225	209	1	1	4.0	0	

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
98	4	0.0	1.0	
103	4	0.0	1.0	
111	4	0.0	2.0	
127	4	0.0	1.0	
128	4	0.0	1.0	
...	
262758	1	1.0	1.0	
262759	1	1.0	1.0	
262760	1	1.0	1.0	
262764	1	1.0	5.0	
262952	1	1.0	2.0	

	Product_Category_2	Product_Category_3	Purchase
98	8.0	17.0	12099.0
103	11.0	16.0	12098.0
111	4.0	15.0	9564.0
127	5.0	9.0	15361.0
128	2.0	15.0	15770.0
...
262758	2.0	15.0	11512.0
262759	15.0	18.0	11521.0
262760	2.0	13.0	3988.0
262764	8.0	17.0	5444.0
262952	3.0	4.0	3240.0

[11904 rows x 12 columns]

sample_size = 1000

sample = new_data.sample(sample_size, random_state=0)

new_data.mean()

User_ID	2525.070733
Product_ID	222.067624
Gender	1.000000
Age	1.000000
Occupation	7.042255
City_Category	1.063172
Stay_In_Current_City_Years	1.866683
Marital_Status	0.194892
Product_Category_1	2.484207
Product_Category_2	6.584677
Product_Category_3	12.471522
Purchase	11785.188172

dtype: float64

new_data.Purchase.std()

5080.322126868783

pos_mean=11785

#one sample t-test

from scipy.stats import ttest_1samp

t_stat, p_value = ttest_1samp(sample['Purchase'], pos_mean)

print(t_stat, p_value)

-0.21540497598886765 0.8294955587149927


```
# P-value is less than 0.05, reject the null hypothesis.
# Null Hypothesis will be accepted
# therefore, the mean purchase for men aged 18-25 is 10000.
```

B) It was observed that the percentage of women of the age that spend more than 10000 was 35%. Is it still the same?

- One Sample Test for Proportion

```
#null hypothesis - proportion is 35%.  
#alternate hypothesis - proportion is not 35%.  
  
new_data = df.loc[(df['Purchase'] > 10000)]  
  
new_data.mean()  
  
User_ID                2777.960883  
Product_ID            221.135756  
Gender                 0.793987  
Age                   2.459020  
Occupation             8.283276  
City_Category          1.125491  
Stay_In_Current_City_Years 1.861065  
Marital_Status         0.403786  
Product_Category_1     1.808018  
Product_Category_2     6.055254  
Product_Category_3    12.360397  
Purchase              14982.258210  
dtype: float64  
  
new_data.Purchase.std()  
  
3040.3040968756654  
  
#hypothesised value  
po=0.35  
  
#number of observations  
  
nobs = len(new_data['Gender'])  
  
#number of women observations  
count=new_data['Gender'].value_counts()[0]  
  
from statsmodels.stats.proportion import proportions_ztest  
z_stat, p_val = proportions_ztest(count=count,  
                                   nobs=nobs,  
                                   value=po,  
                                   alternative="two-sided",  
                                   prop_var=False)  
  
print(z_stat, p_val)
```

```
-79.12020590883206 0.0
```

```
#p-value is less than 0.05, reject the null hypothesis.  
#the proportion of women spending more than 10000, is not 35%.
```

c. Is the average purchase made by men and women of the age 18-25 same?

- Two Sample test for Means

```
#null hypothesis - average spends are equal  
#alternate hypothesis - average spends are not equal  
men = df.loc[(df['Gender'] == 1)& (df['Age'] == 1)]  
women = df.loc[(df['Gender'] == 0) & (df['Age'] == 1)]  
  
#creating samples  
data_men_sample = men.sample(500, random_state=0)  
data_women_sample = women.sample(500, random_state=0)  
  
#checking variances of the two samples  
print(data_men_sample.Purchase.var())  
print(data_women_sample.Purchase.var())  
  
25403579.49849695  
26680870.93292181  
  
#two sample t-test for unequal variances  
from scipy.stats import ttest_ind  
  
t_stat_2, p_val_2 = ttest_ind(data_men_sample.Purchase,  
data_women_sample.Purchase, equal_var=False)  
print(t_stat_2, p_val_2)  
  
3.4922719108842966 0.0004999285373589167
```

```
#we can reject the null hypothesis using the test statistic and since p-value is less than 0.05.
```

```
#the average purchases are not the same.
```

D. Is the percentage of men who have spend more than 10000 same for the ages 18-25 and 26-35

- Two Sample test for Proportion

```

data_age1 = df.loc[(df['Age'] == 1) & (df['Purchase'] > 10000)]
data_age2 = df.loc[(df['Age'] == 2) & (df['Purchase'] > 10000)]

data_age1_sample = data_age1.sample(1000, random_state=0)
data_age2_sample = data_age2.sample(1000, random_state=0)

count = [(data_age1_sample['Gender'] == 1).sum(),
          (data_age2_sample['Gender'] == 1).sum()]

nobs = [(len(data_age1_sample)), len(data_age2_sample)]

from statsmodels.stats.proportion import proportions_ztest
stat_2sample, p_value_2sample = proportions_ztest(count=count,
                                                  nobs=nobs,
                                                  value=0,
                                                  alternative='two-
sided',
                                                  prop_var=False)

print(stat_2sample, p_value_2sample)
-0.759111307093946 0.44778597581119517

```

#p value is more than 0.05, cannot reject the null hypothesis. Null hypothesis is accepted

#therefore, Percentage of the men in both the age groups who have spent more than 10000 is same